Port-Au-Prince Calling: Social Network Response to Social Unrest using Mobile Phone Metadata in Haiti[*]

Daniel S. Putman[†]   Travis J. Lybbert[‡]

This Version: September 7, 2023
Click Here For Most Current Version

Abstract

Which connections do people value most in times of crisis? We examine a period of social unrest in Haiti to understand how social networks respond to social unrest. We construct communication networks using mobile phone metadata from a major mobile network operator in Haiti and a detailed geo-referenced timeline of severe unrest. These episodes are geographically isolated, persist for a matter of days, and vary in their degree of coordination and predictability. We use the more spontaneous of these events to estimate how calling behavior responds day-to-day. Estimating treatment effects using a difference-in-differences estimator robust to variation in treatment timing with heterogeneous treatment effects, we find that daily contacts decrease but total talk time remains constant. These results are consistent with predictions from a theoretical model of social network response to social unrest. However, the number of daily contacts with network neighbors who are strong ties and have low degree does not fall during the period of social unrest, in contrast to other subgroups. This finding suggests a pattern of checking in on close friends, family or other associates in lieu of broader information search.

Keywords: Social Networks, Social Unrest, Digital Trace Data, Call Detail Records, Mobile Phone, Haiti

JEL Codes: D83, L14, O12, O17, Z13

---

1 Introduction

"When you are in trouble, you find out who your friends are." -Haitian proverb[1]

This oft-cited saying in Haiti reflects the activation of social networks in response to crises. While friends, relatives, or neighbors might provide access to cash, credit, or other material aid through risk-sharing networks, social infrastructure does not end with such transfers; networks can also play the vital role of providing information as well as emotional and psychological support in these instances (Aldrich and Meyer, 2015; Blumenstock et al., 2016; Metaxa-Kakavouli et al., 2018). In Haiti specifically, social networks play an important role in accessing resources and information in the wake of crisis (Rahill et al., 2014; Sommerfeldt, 2015). Such *covariate shocks*—shocks correlated across households within a geographic area—can come in many forms. While natural disasters like earthquakes or hurricanes can precipitate intense times of "trouble" across a wide swath of a population, localized disruption due to protests, looting and civil unrest can be equally troubling albeit in a more isolated manner. With a focus on the latter, we aim to understand how troubling times trigger differential social network usage by altering communication incentives.

Haitians have had ample opportunities to find out who their friends are, as this saying suggests. In 2010, a 7.0 Earthquake struck near the capital of Port-Au-Prince, causing massive loss of life, damage to infrastructure, and human displacement (Bengtsson et al., 2011; Rahill et al., 2014). More recently, Haiti has been rattled by a steady stream of civil strife, only some of which is covered by international news media. Starting in 2018, Haiti faced a period of social unrest in response to reports of government corruption and embezzlement of loans from Venezuela's PetroCaribe program.[2] This period included chaotic and decentralized mass demonstrations with crowd sizes of up to several thousands, and roadblocks restricting mobility (Charles, 2019). While these acute unrest episodes inflicted far less permanent damage to physical infrastructure than natural disasters, they were no less disruptive in their immediate vicinity. The demonstrations often turned violent, including cases of tire burning, shootings, and rock throwing leading to the cancellation of Carnival events (Associated Press, 2019; Lemaire, 2019; HaitiLibre, 2019). As compared to natural disasters, they constitute a cleaner shock to the demand for information because they are not confounded by collateral damage to communication infrastructure.

In this paper, we study how communication networks respond to—rather than mobilize—acute civil unrest by considering how these localized shocks shift incentives to communicate throughout these networks. Social and political movements, including protests and civil unrest, have always leveraged social networks in order to motivate and mobilize individuals (Campbell, 2013; Steinert-Threlkeld et al., 2015; Bursztyn et al., 2021; Eubank and Kronick, 2021). However, social networks are not only connected to social unrest through political mobilization. In particular, networks play a key role in the diffusion of information in uncertain and changing states of the world (Conley and Udry, 2010; Banerjee et al., 2013; Magnan et al., 2015; Banerjee et al., 2018, 2019b; Beaman et al., 2021). In the case of social unrest, as incentives shift in

---

[1]In Hatian Creole: "Se lè ou nan male w'ap konn kiyes moun ki zanmi w" (Déralciné and Jackson, 2015). Another iteration goes "se lè ou nan bezwen, ou konn ki moun ki zanmi ou" or "when we are in need, we know who our friends are." (Creole Haiti).
[2]News of this embezzlement was first broken to the public in 2017 by a probe by Haiti's senate but did not boil over into protests until mid-2018 (Charles, 2017, 2019).

times of stress and disruption, individuals activate their social networks, connecting and communicating with links new and old. We aim to characterize a broader network that includes bystanders whose daily lives are disrupted by localized unrest, rather than focusing on communication among protesters and organizers. This broader inquiry evaluates how such bystanders tap into their networks to cope with this greater uncertainty, and the risk to life and property. In doing so, we leverage the digital trace data left behind by mobile phone calls made during these episodes of social unrest.

One of the oft-cited benefits of Information and Communication Technology (ICT) for vulnerable populations is the added resilience that can come from better and more timely information, especially in advance of or immediately after major shocks. We aim to understand how those in close proximity to localized unrest use ICT-mediated social networks to cope with these shocks and the implications of this response for their networks. The informal spread of information about social unrest could similarly serve to mitigate the disruption caused by social unrest. Additionally, calling behavior during crisis is revealing in that it expresses who and what people value. Do those who face unrest in their neighborhoods turn outward for support and information, or do they restrict their communications to their close friends, family and associates?

A few prior studies document how ICT-mediated social networks are activated in respond to major shocks. Blumenstock et al. (2016) leverage an earthquake and document the flow of assistance (in the form of airtime transfers) into the affected area of Rwanda through reciprocal ties mobile communication networks. Jia et al. (2021) assess mobile communication patterns in the wake of the 2013 Ya'an earthquake in China and find that people are more likely to call family members when they have more common friends outside of their families. Romero et al. (2016) examine the effects of asset price shocks on the communication network between employees at a hedge fund, finding that information flowed less freely within the organization during these shocks. These studies feature settings where major covariate shocks unexpectedly affect entire samples or populations. While these are devastating and costly shocks, they are blunt, one-off events that provide little geographic variation with which to identify social network responses. For that, more frequent and localized shocks are more useful.[3] We leverage acute and spatially-isolated spells of social unrest in Haiti to understand how uncertainty driven by social unrest activates ICT-mediated social networks. That is, we estimate the impact of these episodes of social unrest on social networks.

Our empirical strategy is three-fold. First, we examine a spell of social unrest in Haiti in early 2019, leveraging acute and spatially isolated unrest as exogenous shocks to the value of mobile communication. In particular, we restrict our focus to January and February in the Port-Au-Prince area, a time and place that experienced a marked escalation of social unrest. We argue that these spatially isolated shocks serve as a natural experiment suitable for a modified difference-in-differences empirical strategy. Second, to limit issues with anticipation effects, we restrict our analysis to the least predictable events, excluding protests from our definition of treatment. Third, we estimate our treatment effects using the DID^M estimator presented in de Chaisemartin and D'Haultfœuille (2020), which is robust to variation in treatment timing

---

[3]Blumenstock et al. (2020) uses a similar kind of variation, but aims to understand how firms' location decisions are shaped by violence in Afghanistan.

with heterogeneous treatment effects, and flexible enough to allow units to leave treatment.[4] This estimator yields instantaneous treatment effects as areas of Port-au-Prince fall into and out of spells of social unrest.[5] Using this empirical strategy, we estimate the network response using mobile phone metadata from a large telecommunications provider in Haiti. These ICT-mediated social networks yields census-like coverage without the prohibitive expense of in-person surveys. Moreover, the fine grained temporal nature of these communication networks yields the possibility to look at short run changes in network usage in response to shocks. We aggregate network activity into a daily panel centered around this spell of social unrest. Then, we assign users into treatment if events of social unrest occur in close proximity to their most used tower.

To frame this empirical analysis, we first construct a theoretical model of network response to social unrest, building on work by Björkegren (2019). In this model, the effect on calling behavior depends on the underlying strength of social connection—an important feature in our context—peer information, and the cost of calling—which we define as the price of airtime plus (what we term) an attention cost (Rahill et al., 2014).[6] We categorize ties as either unrecognized, recognized, or trusted, in order of increasing strength. The model predicts that while unrecognized ties are never contacted, merely recognized ties cease to be contacted during episodes of social unrest. On the other hand, trusted ties are contacted. The intensity of connection behaves according to a combination of their strength of tie, and if they are likely to have information about the unrest. Using these predictions, we hypothesize that social unrest will increase the duration of calls to the strongest (trusted) ties as a result of their accompanying information shock, reduce communication with unrecognized ties and the weakest trusted ties, and increase or decrease communication to other trusted ties based on their information about unrest events. Notably, this should reduce number of contacts on average.

Using a five week window, using DID^M we estimate that non-protest social unrest reduces contacts by 0.05 per day and total calls by 0.1 per day. However, total duration (i.e., time spent talking) stays roughly constant.[7] This pattern of results might suggests people talk with a smaller set of important contacts – consistent with the hypotheses of the theoretical model described above. To explore who people value in a time of crisis, we decompose these effects by whether contacts are strong ties or weak ties at baseline. Similarly, we decompose results by whether contacts have high or low degree at baseline, which we view as proxying for peer information. When we restrict to only low degree strong ties, there is no reduction in the number of contacts–in contrast the main result and the other subgroups. This suggest that people most value checking in on their closest friends, family, or associates as the uncertainty of unrest arises. Interestingly, this story is also consistent with evidence from disparate shocks such as earthquakes and

---

[4]In particular, in such contexts, two-way fixed effects (TWFE) estimators may be biased due to issues with weighting of group specific treatment effects (Goodman-Bacon, 2019; de Chaisemartin and D'Haultfœuille, 2020). By weighting, we eliminate both those comparisons are included in error (i.e., that should have zero weights), and the issues of negative and arbitrary weights that occur as an artifact of TWFE.

[5]By instantaneous, we mean the treatment effect using the period just before and just after treatment.

[6]Essentially, this is an opportunity cost from diverting limited attention away from monitoring one's surroundings.

[7]DID^M is our preferred estimator. We diagnose differences between TWFE and DID^M where these differences should be attributable to differences in weighting of heterogeneous treatment effects in a context with variation in treatment timing. We find that the weighting provided by TWFE yields substantial differences in treatment effects, consistent with tests introduced by de Chaisemartin and D'Haultfœuille (2020).

stock crashes (Romero et al., 2016; Blumenstock et al., 2016; Jia et al., 2021). In contrast to past work, we are able to draw on spatial variation in order to provide evidence using a difference-in-differences framework and are able to avoid the impact on infrastructure often inherent in natural disasters.

Our results matter for policymakers who are facing crises such as natural disasters or sit within organizations that are exposed to financial shocks. While we hypothesize that major shocks might spur contact to high degree peers who might have better access to information about what is happening on the ground, our results are inconsistent with this kind of information search behavior. Such behavior might allow networks to make use of 'small worlds' properties to efficiently diffuse information. However, this might require people to reach out to central ties even when they are not as close of connections (Granovetter, 1973; Strogatz and Watts, 1998; Centola et al., 2007). Instead, we see that actors within these networks form silos, speaking less to people outside their close ties. Our results are also related to the finding that social networks are optimized for reciprocal support (i.e., the exchange of favors) in contrast to the flow of information (Jackson et al., 2012; Blumenstock et al., 2019). In light of this, we hypothesize there is a double penalty in information diffusion in the wake of crisis: not only are network structures oriented toward support as opposed to the flow of information, but also people do not value connecting with those in their network who might engender information flow. Where different populations are not equally well connected, outside of broad-based information campaigns, this could lead to cases some groups are much less informed about the risks posed by a crisis.[8] Such issues with informal information spread points to the importance of such information campaigns which may use one-way SMS blasts, diesseminate information through crowd-sourcing platforms, or harness individuals' movements or activities as relayed by their mobile devices (Azid et al., 2015; See, 2019; Lu et al., 2012; Sakaki et al., 2010). This suggests that while there may be some value of private communication, a robust disaster response should not rely on word-of-mouth alone.[9]

2 Background

2.1 Related Literature

2.1.1 Social, Economic, and Digital Networks

As access to administrative data has grown, digital trace data has become an attractive approach to understanding social and economic network structure. This data is highly detailed and maps intuitively onto the data structures used in social network analysis. Calls, SMS, and mobile money products tend to perform similar functions to their in-person counterparts, the degree to which has been documented in literature measuring online social networks and comparing these to offline social networks. Using both social media data and survey data, Dunbar et al. (2015) finds that online and in-person networks have similar structures,

---

[8]For a similar example from a public health crisis, Bennett et al. (2015) document evidence of learning in response to SARS in Taiwan where recent arrivals miss out on learning about the novel virus.

[9]Some past work points to the importance of social capital in disaster response (e.g., Aldrich and Meyer, 2015). This doesn't necessarily provide evidence to the contrary. Indeed, it may be instructive of issues that arise when social capital is weak.

including similar frequency distributions of contacts.[10] Similarly Bisbee and Larson (2017) presents survey data about online and offline ties and finds strong similarities in the nature of online and offline relationships, even as identities of ties differ. In the context of favor exchange networks, Blumenstock et al. (2019) finds a preference for interconnected mobile networks among migrants, where relationships are supported by common contacts. This echoes results from in-person social networks constructed around support and favor exchange (Jackson et al., 2012). Similarly, where in-person risk sharing networks help cope with idiosyncratic shocks (Fafchamps and Lund, 2003; De Weerdt and Dercon, 2006; Ambrus et al., 2014), mobile money facilitates risk sharing in the wake of both idiosyncratic and (locally) covariate shocks (Jack and Suri, 2014; Blumenstock et al., 2016; Riley, 2018).

2.1.2 ICT Mediated Social Networks and Resiliency in the Wake of Crises

Work documenting the use of networks in resilience to shocks tends to focus on risk-sharing networks. Despite the relative scarcity of quantitative evidence about word-of-mouth information during crises, it is often cited as a vital element of disaster risk reduction (Aldrich and Meyer, 2015; Bernier and Meinzen-Dick, 2020). Some recent work documents the search for information using communication networks in the wake of crises. Jia et al. (2021) assesses mobile communication patterns in the wake of the 2013 Ya'an earthquake in China, paying special attention to how the interconnectedness of family networks affects the dynamics of responses within families. They find that those families (defined by those on the same family plan) who have more out of family contacts in common are more likely to contact each other shortly after the earthquake, as opposed to other contacts outside of the family structure. Romero et al. (2016) examines the effects of price shocks on the communication network between employees at a hedge fund. During large stock price drops, the data shows increased siloing among closely connected members of the firm. They argue that response, affectionately described as 'turtling up,' represents a reduction in the ability of the organization to gather information during such price shocks. Additionally, some work also documents how word-of-mouth information can change behavioral response to crises. For example, Bennett et al. (2015) documents evidence of learning in response to SARS in Taiwan by comparing the behavior of recent migrants to others in the area. They interpret differences in preventative actions in response to local (as opposed to national) information as resulting from social interactions with peers.

In addition to their use in understanding behavior, there is a growing use of CDR data to detect anomalies in usage patterns and mobility as a form of real-time disaster monitoring. While the goal of this 'user-as-sensor' approach is not to understand behavior *per se*, the features which warn of crises are instructive. For example, natural disasters tend to drive increases in mobile phone usage and mobility (Bengtsson et al., 2011; Dobra et al., 2015; Flowminder, 2021). On the other hand, violence against civilians and protests have less consistent signatures. Anomalies sometimes include reductions in mobile phone usage and mobility while in other cases, there were increases in both usage and mobility the day following the event (Dobra et al., 2015; Gundogdu et al., 2016). Such differences in response may prove important in interpreting our

---

[10]This includes Twitter and Facebook networks have similar contact frequency distributions as in-person contacts. That is, for a given intensity of interaction, on average users interact with about the same number of people they might in an offline setting as measured by survey.

results.

2.1.3 ICTs and the Mobilization of Social Unrest

While we focus on bystanders, social networks play a well documented role in mobilizing political participation (Campbell, 2013). Protest mobilization in the United States civil rights era took place in movement organizations (i.e., the formal or informal core of activists) as opposed to more prosaic word-of-mouth (Andrews and Biggs, 2006). While information might flow more easily via word-of-mouth, this might be explained by the acute risks involved in the decisions to protest (as it was during this era). This contrasts with what appear to be more spontaneous forms of social unrest which arose during the Arab Spring. Steinert-Threlkeld et al. (2015) documents social media activity in 16 countries during the Arab Spring, finding that decentralized social media activity correlates with the coordination of protests shortly thereafter. However, despite modern communication platforms, evidence points to similarities with the movement organizations. In particular, modern mobilization seems to take place within tightly knit portions of networks (Bursztyn et al., 2021), and may be bolstered by visibility and strong moral rhetoric among a set of likely minded people (Mooijman et al., 2018; Eubank and Kronick, 2021).

2.2 Data and Context

2.2.1 Social Networks and Resiliency in Haiti

Past research from an international context stresses the role of informal risk sharing arrangements in coping with risk (Fafchamps and Lund, 2003; De Weerdt and Dercon, 2006). Such reciprocal arrangements are also a part of life in Haiti, growing out of systems of *lakou* in post-colonial Haiti (Miller, 2013; Merilus, 2015). The *lakou*, which translates to courtyard, is a land management system made of five to seven families on a shared plot of land. Due to Haiti's past as a slave colony, the *lakou* contained those who were tightly knit to the inheriting family, including both kin and 'fictive' kin. The *lakou* was also home to *konbit*, a reciprocal agricultural labor sharing arrangement. While the *lakou* is a feature of rural life, Miller (2013) suggests that the structure of the Lakou is reflected in post-disaster settlements of urban Port-au-Prince after the 2010 earthquake.

As we consider resilience within neighborhoods and camps in Port-Au-Prince in the wake of the 2010 earthquake, it clear that networks and other forms of social capital play a role in resilience. Sommerfeldt (2015) documents information seeking repertoires in refugee camps, finding word-of-mouth to be the most commonly and intensively used. Rahill et al. (2014) map out four concepts of social capital that guide diffusion of information and resources in camps. In particular, relationships can take on three categories of increasing closeness. The first level of closeness is *rekonnèt*, or the recognition of each other from positive previous interaction. Building on repeated positive interactions, *konfyans* can be though of as trust. *Moun pa* indicates membership in each others' inner circles, a step beyond trust. Finally, *pati pri* indicates the decision to take sides with that person over another. Rahill et al. (2014) show that in the resource scarce context of post-disaster Haiti, *pati pri* is seen as increasing pre-existing inequalities. Even in cases of information sharing, trusted ties may be important due to the combination of anxiety and information
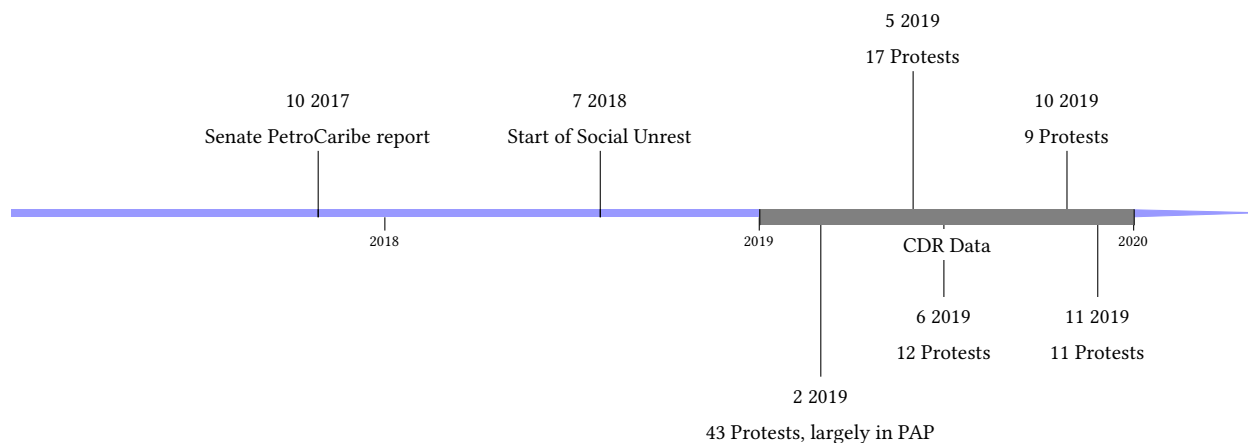
Figure 1: Highlights of Protest Events in Haiti, 2017-2019

uncertainty. Oh et al. (2010) documents the diffusion of rumors and misinformation in this period. Finally, broader social cohesion within neighborhoods of Port-Au-Prince is associated with the self-assessment of community resilience (Patel and Gleason, 2018).

2.2.2   Social Unrest in Haiti

To measure spells of social unrest, we use a timeline of social unrest events throughout 2019, taking into account kinds of unrest on various days throughout the year. Starting in 2018, Haiti faced a period of social unrest in response to reports of government corruption and embezzlement of loans from Venezuela's PetroCaribe program (Charles, 2017). The high-level view of this timeline can be seen in Figure 1. While it is difficult to estimates the total participation in the protests, these spells of unrest were very disruptive. For example, they forced the cancellation of Carnival events in 2019 and 2020. On February 17th, Prime Minister Jean Henry Céant noted that "it's been ten days since children have been unable to go to school, hospitals can't provide healthcare, big businesses and small businesses can't function." Similarly, citizens noted in dismay "people can not eat," "the economy is bad, really bad," "when you are hungry, it's not sweet" (Charles, 2019). Protests were not just large and intense, but also likely salient to bystanders as insecurity is a key issue identified by Haitians. According to the 2021 AmericasBarometer survey, 45% of Haitians placed insecurity as the most important problem they faced (Gélineau et al., 2021). In 2021, 66% of Haitians felt insecure in their neighborhoods.[11]

Using this timeline, we zoom in on those events taking place in January-February, 2019. The timeline itself draws on embassy security reports, newspaper articles, and social media activity to determine when

---

[11]While this rose over the time period in question, 55% of Haitians reported feeling insecure in their neighborhood in 2017. Notably, the 2021 survey data was collected prior to the assassination of Haitian President Jovenal Moise.
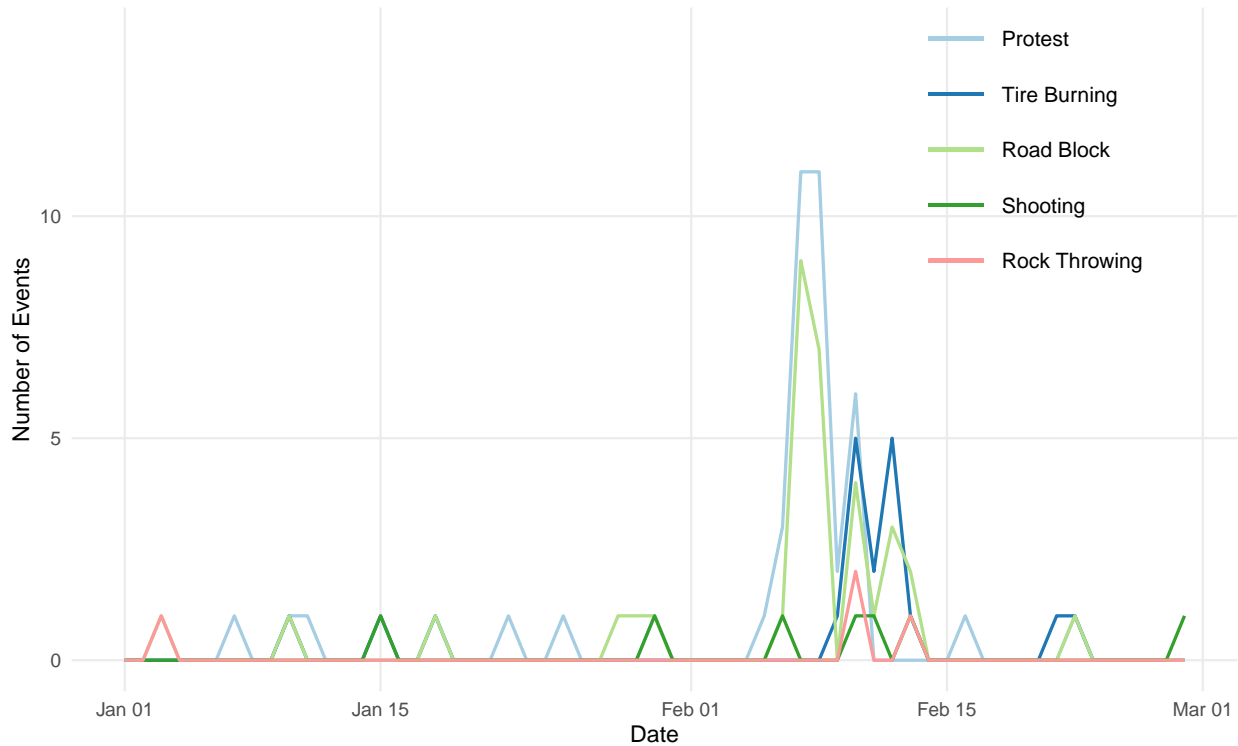
Figure 2: Number of Social Unrest Events by Day in Port-Au-Prince, January and February 2019

and where social unrest took place (Pierre-Charles et al., 2020). The data contains information about the type of social unrest, geographic location, start time, and level of severity of the unrest (as determined by the US embassy).[12] The types of social unrest identified are protests, roadblocks, shootings, rock throwing, and tire burning, which covers the large majority of events throughout the year.[13] We plot these five forms of social unrest by their frequency in January and February in Figure 2. Notably, the increase in number of events in early February changed the codes used by US embassy-provided security updates. While the earlier unrest was coded "avoid area," from the 7th to the 10th of February, the code was upgraded to "home restriction." Finally, events on the 11th to the 21st carried a "shelter in place" code before returning to "avoid area" codes at the end of the month. To get a general sense of the location of unrest, we map events in Port-Au-Prince in Figure 3 as well as the individual types of social unrest in Figures 9-13. Notably, when examining the disaggregated maps, one pattern emerges in the distribution of protests relative to other forms of social unrest: While protests clustered around Champs de Mars (along with the presidential palace and government administrative complexes), other events tended to cluster more to the east, in more residential areas.

These five forms of social unrest differ on a number of margins. One is the scale of coordination

---

[12]In particular, we list the most likely coordinates to each event, and the exact coordinates if possible. When a specific intersection of streets is in mentioned in a report the location is quite precise. When one of the many small neighborhoods are mentioned we list the coordinates of the main intersection in this neighborhood, which is the most likely place for protest to take place. Neighborhoods are relatively small and so even imprecise placement should bewithin 500m of the true location.

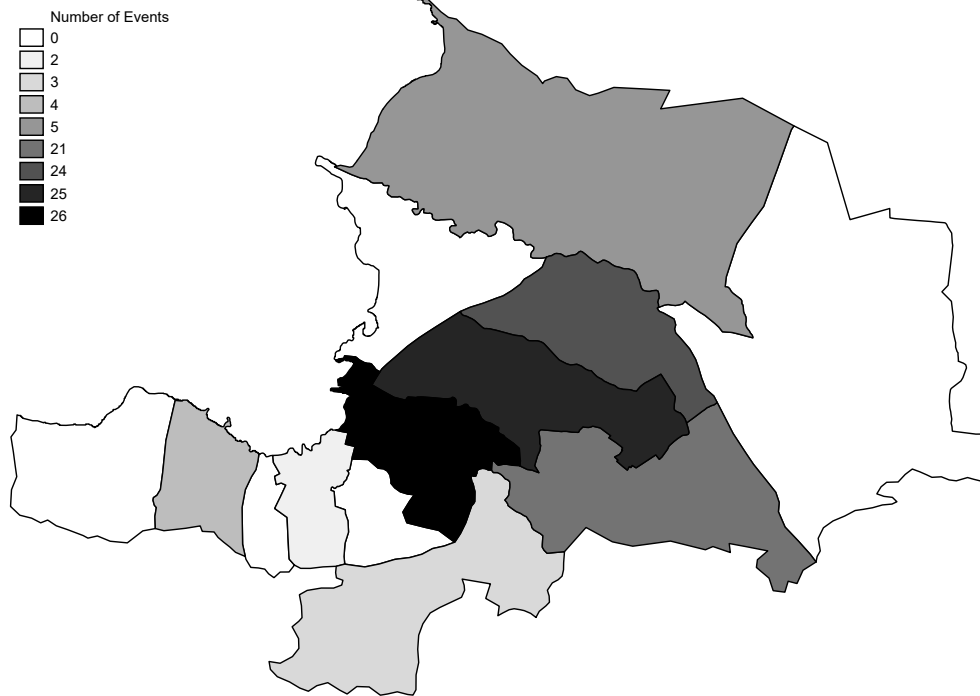[13]The remaining unspecified events are dropped.

Figure 3: Social Unrest Events in Port-au-Prince, January and February 2019.

necessary for their occurrence. Leading the way are protests which require a high degree of coordination, enough to ensure a good turnout of protesters. Second is tire burning, which are reported to require around five experienced participants.[14] This is followed by roadblocks, which require just enough people to move a car or some other impediment.[15] Finally, shootings and rock throwing require the lowest level of coordination.

Other margins we can consider include the possibility of anticipation and how geographically local they are. Notably, while protests may be most disruptive, they are also possible to anticipate due to their high degree of coordination. The other events are less easy to anticipate, and some (shooting and rock throwing) are considerably more spontaneous than others. Likewise, it's important to consider the geographic scale. It's intuitive that both protests and roadblocks have implications that extend beyond their specific location as they often take place on main thoroughfares and thus have the ability to jam up the city. In contrast, shooting, tire burning, and rock throwing may be the most acute in localities.

2.2.3 Mobile Phone Metadata

The mobile phone metadata used in this project comes from the largest mobile network operator in Haiti as part of a long-term research collaboration that gives us access to these data. This mobile network operator

---

[14]"How many people does it take to burn a tire?" is a question with an answer: typically five. Two people carry the tire, two carry gasoline or some other inflammatory substances, and the last one lights the tire with a match (or lighter). See Anonymous (2015) for more.

[15]This could be more than five people, but needs not be, so the minimum to create such a roadblock could happen with lower coordination. Of course, it could also feature more coordination in specific cases.

Table 1: Baseline Outcomes

Statistic	Mean	SD	Median
Degree	18.2	26.3	11
Calls	88.5	135.0	38
Duration (seconds)	10815.8	20463.9	3038

has a dominant market share in the Haitian telecommunications market, which reduces concerns about sampled networks (Chandrasekhar and Lewis, 2016). Metadata like this is often referred to as Call Detail Records (CDRs) and, in our case, includes transaction-level records of calls and text messages: caller ID, recipient ID, date time, duration of calls in seconds, caller tower, recipient tower (not included for SMS or calls out of network), and traffic type (voice or SMS). Under the auspices of our research agreement, we have access to multiple years of CDRs from this network operator, but in this analysis, we use data from 2019 and use communication patterns as reflected in these data as the basis for constructing latent social networks.

2.2.4 Call Outcomes

To construct outcomes to test these hypotheses, we focus on individual cellphone usage. These outcomes included the number of unique contacts, total calls, and total duration of calls, each of which has an ana-logue in social network centralities. First, we compute degree centrality, yielding the number of unique contacts in a day for all nodes. Second, we can compute "weighted" degree centralities which yield both the total number of calls each day and the total duration of those calls. We compute a baseline network of voice calls using data from the first three weeks of the year. Table 1 presents these statistics at baseline.[16]

2.2.5 Baseline Network: Strong Ties and Centrality

To measure the strength of network connections prior to the onset of the 2019 unrest episodes, we leverage communication patterns as represented in the CDRs during this baseline period. Specifically, we use the total duration of calls between individuals in the baseline network to define strong ties and the number of contacts at baseline to define central nodes, which should be well informed. We define strong ties as those above the 80th percentile in call duration (after dropping dyads with no call duration between them), so about one in five individuals in one's network neighborhood serve as strong ties. We define a high-degree node as one that has more than 18 contacts at baseline, the average number of contacts in this data. These nodes are above the 67th percentile in the distribution, making up a third of the network in this definition. For more on these the process we used to reach these definitions, see Appendix B.1.

---

[16]We also compute the global clustering coefficient of the graph to measure how tightly knit the graph is as a whole. More specifically, the clustering coefficient answers the question: if i is connected to j, who is also connected to k, what is the probability that i also connected to k? In the baseline network, we find a clustering coefficient of 2.7%, suggesting many "friends of friends" are unknown to the ego.

3 Theoretical Model

We construct a model of network response that builds on the model presented in Björkegren (2019). In particular, we model the response of those who are exposed to social unrest by geographic proximity. Our adaption differs in a number of ways from the original model. First, while the model is used in that work to understand calling behavior as a motivation for handset adoption, we find it useful to model the calling decision as an end unto itself. Second, the model there provides structure for the empirical exercise. Here, we use the model as a vehicle to generate hypotheses but do not extend it to provide structure for our estimates. Still, we find some features of the model very useful. In particular, the utility gained from making calls accords with eight reasonable properties related to demand for calling, at least six of which remain relevant in this context. Therefore, we adapt the utility function, embedding within it a model of expectations about the informational value of calling specific network neighbors during crisis.

3.1 Social Network Response to Social Unrest

3.1.1 Setup: Utility and Cost Functions

We specify utility for the value of calls with the following functional form, adapting Björkegren (2019):

$$v_{ij}(d_{ijt}, \epsilon_{ijt}) = \alpha_{ij}d_{ijt} - \frac{1}{\epsilon_{ijt}}\frac{d_{ijt}^{\gamma}}{\gamma}. \tag{1}$$

where d is duration (in seconds), ϵ is a communication shock, $\gamma > 1$ controls the decline of marginal returns, and α controls the intercept for marginal utility of calling. In line with the qualitative work of Rahill et al. (2014), we allow α to depend on both the caller i and the receiver j, which embodies tie strength. In relative terms, those with high values of α can be thought of as strong ties and those with low values, weak ties. Likewise, the communication shock ϵ plays an important role in the model. In particular, this shock reflects the informational content to i of a call to j, conditional on the social unrest on day t.[17]

We suppose no fixed costs but a marginal cost of calling c_{it}.

$$c_{it} = p + \phi(z_{it}) \tag{2}$$

where p is the per second price of calling and ϕ is an attention cost of calling, which depends on social unrest, z_{it}. The intuition here is that as social unrest strikes, people become more vigilant monitoring their surroundings. Phone calls are a distraction from monitoring one's environment and therefore this reduction in attention adds to the cost of making a call, which we take to be linear in call duration. This assumption comports with both the salience of insecurity during the study period and the well documented attention impairments due to cell phone usage (Gélineau et al., 2021).

The utility and cost functions feature eight reasonable properties for cellular call behavior, six of which are also important in this application. For more on these properties, please see Appendix A.1.

---

[17]While we abstract away from this, in reality, we can think about shocks as a mix of expectations and the actual value of information gained after a call begins. For example, I might expect my friend has a great deal of information about the current unrest, but they may relay to me that they do not after I call them, adjusting down the "shock" and ending my call.

3.1.2 Optimization: Caller's Problem

The caller's problem is as follows:

$$\max_{d_{ijt} \geq 0, \; j \in N(i)} U_i(\mathbf{d}_t) = \sum_{j \in N(i)} \left[\frac{1}{\beta} v_{ij}(d_{ijt}, \epsilon_{ijt}(z_{it})) - (p + \phi(z_{it}))d_{ijt} \right] \tag{3}$$

where $N(i)$ is agent i's neighborhood and β converts units from utils to dollars. In addition, we allow the information shock to depend on social unrest, z_{it}, since social unrest drives the search for information.

3.1.3 Solutions: Call Duration

Setting marginal cost equal to marginal utility, when $d_{ijt} > 0$ we arrive at a solution of

$$d^*_{ijt}(\epsilon_{ijt}(z_{it}), \phi(z_{it})) = [\epsilon_{ijt}(z_{it}) (\alpha_{ij} - \beta (\phi(z_{it}) + p))]^{\frac{1}{\gamma - 1}}. \tag{4}$$

When is $d_{ijt} = 0$? In this model, in cases where any positive duration yields negative utility, no call is made. Since $v_{ij}(d = 0) = 0$, the caller chooses a duration of zero over any positive duration. In particular, when

$$\alpha_{ij} < \beta(\phi(z_{it}) + p) \tag{5}$$

a call will not be made. While we take the neighborhood of i as given, this expression can also be thought of as a way to implicitly define neighborhood.[18] We refer to those who might be called when there is no social unrest as being recognized, an expression used to parallel Rahill et al. (2014).[19]

3.2 Comparative Statics

3.2.1 Recognized but Not Trusted Ties Are Not Contacted During Social Unrest

To construct hypotheses about the response of social networks to social unrest, we examine how calls and call duration change when social unrest is switched on. Let $\Delta\phi(z_{it}) = \phi(z_{it} = 1) - \phi(z_{it} = 0)$. Furthermore, let $\phi(z_{it} = 0) = 0$ as in the absence of social unrest we assume there is no relevant threat to pay attention to. Given $\Delta\phi(z_{it}) > 0$ condition (5) yields a first prediction about the decision to call: attention cost driven by social unrest reduces the set of alters one will talk to when they are proximate to social unrest. We refer to those who no longer place calls after social unrest, or those j where $\alpha_{ij} > \beta p$ but $\alpha_{ij} \leq \beta(\phi(z_{it}) + p)$ as recognized but not trusted ties, again paralleling Rahill et al. (2014).[20] If there is any number of these recognized but not trusted ties in the average network neighborhood, we should expect some reduction in the set of contacts one calls on the day of social unrest. It is also useful here to recall that there is another layer of closeness, beyond trust. In particular, stronger connections may be part of one another's inner circles.[21]

---

[18] If $\alpha_{ij} \leq \beta p$, then no calls will ever be made between i and j regardless of social unrest. Therefore, $j \notin N(i)$.

[19] In particular, we are appealing to the concept of *rekonnèt*, (Rahill et al., 2014).

[20] In this case, we are appealing to *konfyans* (Rahill et al., 2014).

[21] Appealing here to *moun pa* (Rahill et al., 2014).

Table 2: Heterogeneous Responses of Dyadic Calling Behavior to Social Unrest

Regime	Ranges of α_{ij}		$d^*_{ijt}(z_{it}=1)$	$\Delta d^*_{ijt}(z_{it})$
Unrecognized	α_{ij}	$\leq \beta p$	$= 0$	$= 0$
Recognized	$\beta p <\ \alpha_{ij}$	$\leq \beta(\phi_{it}+p)$	$= 0$	< 0
Trusted	$\beta(\phi_{it}+p) <\ \alpha_{ij}$	$< \beta\left(\frac{\phi_{it}+(\phi_{it}+p)\%\Delta\epsilon_{ijt}}{\%\Delta\epsilon_{ijt}}\right)$	> 0	< 0
Activated	$\beta\left(\frac{\phi_{it}+(\phi_{it}+p)\%\Delta\epsilon_{ijt}}{\%\Delta\epsilon_{ijt}}\right) \leq\ \alpha_{ij}$		> 0	> 0

Note: $\%\Delta\epsilon_{ijt} = \frac{\Delta\epsilon_{ijt}(z_{it})}{\epsilon_{ijt}(z_{it}=1)}$ and $\phi_{it} = \phi(z_{it})$.

3.2.2 Other Responses Depend on Tie Strength and Information Shocks

Next, we consider the impact of social unrest on trusted ties, or relationships that would have positive call duration during social unrest:

$$\Delta d^*_{ijt}(\epsilon_{ijt}(z_{it}), \phi(z_{it})) = \left(\frac{1}{\gamma-1}\right) [\epsilon_{ijt}(z_{it})(\alpha_{ij} - \beta(\phi(z_{it})+p))]^{\frac{1}{\gamma-1}-1}$$
$$\times [\Delta\epsilon_{ijt}(z_{it})(\alpha_{ij} - \beta(\phi(z_{it})+p)) - \beta\epsilon_{ijt}(z_{it})\Delta\phi(z_{it})]. \tag{6}$$

We simplify this expression:

$$\Delta d^*_{ijt}(\epsilon_{ijt}(z_{it}), \phi(z_{it})) = d^*_{ijt}(z_{it}) \times \frac{\Delta\epsilon_{ijt}(z_{it})(\alpha_{ij} - \beta(\phi(z_{it})+p)) - \beta\epsilon_{ijt}(z_{it})\phi(z_{it})}{(\gamma-1)[\epsilon_{ijt}(\alpha_{ij} - \beta(\phi(z_{it})+p))]} \tag{7}$$

$d^*_{ijt}(z_{it}) > 0$ implies $(\alpha_{ij} - \beta(\phi(z_{it})+p)) > 0$, and $\gamma > 1$. Therefore, the sign of the derivative depends only on the sign of the expression,

$$\Delta\epsilon_{ijt}(z_{it})(\alpha_{ij} - \beta(\phi(z_{it})+p)) - \beta\epsilon_{ijt}(z_{it})\phi(z_{it}). \tag{8}$$

Call duration increases for those dyads where this expression is positive. Several different manipulations of this expression are useful. First, to build intuition, we express the inequality such that response is positive as a ratio compared to one:

$$\left(\frac{\Delta\epsilon_{ijt}(z_{it})}{\Delta\phi(z_{it})}\right)\left(\frac{\alpha_{ij} - \beta(\phi(z_{it})+p)}{\beta\,\epsilon_{ijt}(z_{it})}\right) > 1. \tag{9}$$

The sign of the effect depends on the ratio of the effect of social unrest on informational content of calls (how much more do we have to learn from person j when the world becomes more uncertain?) to the effect of social unrest on attention cost (how much more of your environment do you miss when chatting on the phone when the world becomes uncertain?). Additionally, the likelihood this condition holds increases in dyads with pre-existing strong ties (α_{ij}), falls in costs to i ($\beta(\phi(z_{it})+p)$), and attenuates in the informational shock ($\epsilon_{ijt}(z_{it})$).

Second, we characterize the bound where the value switches from negative to positive in terms of the strength of α_{ij}. This gives us an expression similar to 5. For a given dyadic shock to attention cost and
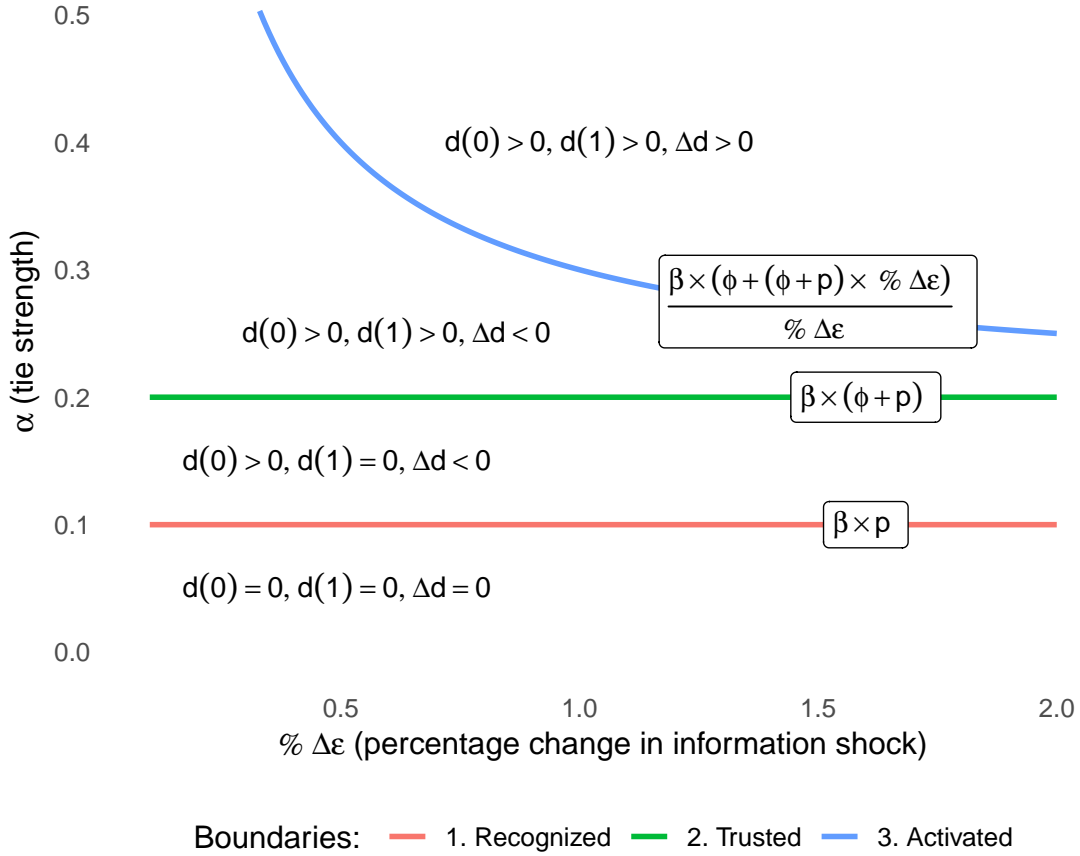
Figure 4: Heterogeneous dyadic calling response by tie strength and change in information shock. $d(0)$ is duration with no social unrest, $d(1)$ is duration with social unrest and Δd is the difference in duration. $\%\Delta\epsilon$ is the percentage change in information shock and α is tie strength. ϕ is attention cost and p is price. Visualization parameters: $\phi = p = 0.1$ and $\beta = 1$.

information, if

$$\alpha_{ij} > \beta \left(\frac{\phi(z_{it} = 1) + (\phi(z_{it} = 1) + p)\frac{\Delta\epsilon_{ijt}(z_{it})}{\epsilon_{ijt}(z_{it})}}{\frac{\Delta\epsilon_{ijt}(z_{it})}{\epsilon_{ijt}(z_{it})}} \right) \tag{10}$$

then we will see an increase in call duration. We refer to ties where equation (10) holds as activated ties.

Duration response is summarized in Table 2. Note that as the percentage change in information shock grows, the right hand expression in inequality 10 converges to $p + \phi(z_{it})$, as is plotted in Figure 4. This suggests that when the information shock is sufficiently large, even relatively less strong ties will almost always be contacted more, given that call duration is positive. Likewise, when the percentage increase in information is sufficiently small, ties will need to be increasingly large to allow for an increase in call duration.

3.2.3 Aggregation of Dyadic Effects to the Network Neighborhood

Translating to neighborhood level calling behavior, the impact on total is the sum of the impacts on dyadic call duration:

$$
\frac{\partial}{\partial z_{it}} \left\{ \sum_{j \in N(i)} d_{ijt}^*(\epsilon_{ijt}, p) \right\} =
$$
$$
\sum_{j \in N(i)} d_{ijt}^* \times \frac{\epsilon_{ijt}'(z_{it})\left(\alpha_{ij} - \beta\left(\phi(z_{it}) + p\right)\right) - \beta\epsilon_{ijt}(z_{it})(\phi'(z_{it}))}{(\gamma - 1)[\epsilon_{ijt}(z_{it})(\alpha_{ij} - \beta(\phi(z_{it}) + p))]} \tag{11}
$$

However, because these impacts depend on j, we cannot make a prediction about the effect on neighborhood level calling behavior without some idea about the dyad level parameters α_{ij} and shocks ϵ_{ijt}.

Considering the number of contacts who are called, the more weak ties there are in the neighborhood, the more the number of contacts will fall. In particular, this reduction in contacts would be the case in a situation where for some households, condition 5 switches on during social unrest. We presume there are some weak tie individuals in an average network neighborhood, so we would expect a reduction on average in the number of contacts made during social unrest.

However, even when there is a reduction in the number of contacts called, there could easily be an increase in total duration, ifwe find that condition 8 holds for sufficient number of other nodes.[22] Given that this depends on the composition of parameters in the neighborhood, we need to proxy for these parameters to build predictions about total duration response. While α should be closely related to past duration of calls, ϵ has much less structure. Therefore, to make informed predictions about total duration we need to explore heterogeneity that might augment the information shock along a given dyad.

3.3 Network Structure and Information Shocks

While we have taken network structure into account when considering strong and weak ties, based on the model we have derived, we have left information shocks as exogenous and without structure. To better guide our hypotheses, we provide more discussion about factors that likely shape how information flows across the networks represented in our CDRs and how information shocks may be differentially transmitted across this network. Three factors in particular are noteworthy and described in detail in this section. First, edges where the alter is more central in this network from the perspective of information diffusion will experience more significant information shocks in the wake of social unrest than those with lesser centrality in the network. Second, edges where the alter's information is more differentiated from the ego's may have higher information shocks. This might be the case when the two individuals share few common friends and therefore end up hearing different bits of information from different sources. Third, individuals whose daily lives and routines bring them into close proximity to localized social unrest likely

---

[22]One useful piece of intuition might come from the sixth feature of the utility function. In particular, changing the marginal cost of a call affects longer calls more than shorter calls. This would suggest that attention cost will yield more resistance against those who already call often, i.e., those with strong ties and or high shocks. This would suggest more of the impact might come from previously low duration callers where information has grown.

know more about what is happening on the ground and thus experience a greater information shock as a result of this unrest.

3.3.1 Diffusion Centrality

Following Banerjee et al. (2019a), we posit that even while individuals do not know the full structure of their surrounding network, they have strong priors about who in their neighborhood is informed, based on who they hear information and gossip from (ultimately) in previous cases. Thus, we posit that ϵ_{ijt} depends specifically on expectations of neighbors' ability to aggregate information. The process by which individuals would identify potentially informed sources brings to mind the "hearing matrix" presented in Banerjee et al. (2019a), which considers the concept of diffusion centrality presented in Banerjee et al. (2013) from a receiver's perspective. However, since we have shown that very weak ties will not be called, the ego will not call out of its network neighborhood. In this case, Banerjee et al. (2019a) shows that diffusion centrality will be proportional to (out)degree centrality. Therefore, $E(\epsilon_{ijt})$ should be directly proportional to alter degree—meaning individuals call their high degree neighbors in a crisis.

3.3.2 The Strength of Weak Ties?

In addition to the ability to aggregate information, it is reasonable to consider the importance of weak ties in the diffusion of information. While this idea has long been discussed in the sociological literature, it has seen development since as bridging social capital (Granovetter, 1973; Aldrich and Meyer, 2015). This idea suggests that those ties you communicate with less frequently are also socially distant, allowing them to diffuse new bits of novel information. In our model, this might be embodied by a inverse relationship between in the strength of tie and the size of ϵ_{ijt}. This is most apparently true for cases where information is diffused through simple contacts, as in the influential study of microfinance by Banerjee et al. (2013). However, when information is in dispute or there is potential for competing stories, the length of these ties may not help for diffusion. In particular, these cases may require complex diffusion where multiple contacts are necessary (Centola and Macy, 2007; Beaman et al., 2021). In fact, we may be in a situation of complex contagion, owing to the potential for rumors and misinformation (Oh et al., 2010).

3.4 Hypotheses

From this model, stylized facts about social networks, and the acuteness of the social unrest shock in this context, we arrive at four hypotheses:

H1: *Social unrest leads to higher network usage among trusted ties.* As social unrest takes place, information shocks become large on average as the surrounding environment becomes unpredictable. Among the neighborhood of trusted ties, these information shocks will lead to increased total duration of calls.

H2: *Merely recognized contacts are not called regardless of information shocks.* People communicate with those they have sufficiently strong ties to but not those weakest (positive) ties, which we term recognized ties. This leads to a reduction in the number of people called during social unrest.

H3: *Trusted ties are activated according to their information shock.*

    (a) Call duration will increase for those who have relatively high information, conditional on their tie strength.

    (b) Call duration will decrease for those who have relatively low information, conditional on their tie strength.

For H3, this is a qualitative expression of the 'sliding scale' condition depicted in equation (10). For example, among ties who are trusted but are not in the inner circle will have higher call duration if the information they hold is useful enough. On the other hand, a member of one's inner circle might be contacted even if they do not have particularly compelling information. We would expect call duration to increase the most among strong ties with high information. Empirically, we posit that the size of the information shock is related to degree centrality.

4 Empirical Strategy

4.1 Difference-in-Differences

We use a difference-in-differences style strategy to estimate the response of social networks to social unrest. The empirical strategy of this paper is based on geographic variation in social unrest. This strategy does not rely on random assignment of social unrest, but rather on the assumptions of parallel trends.

    We start with traditional difference-in-differences approaches to build intuition. The estimating equation for difference-in-differences without individual fixed effects is written

$$y_{ilt} = \alpha + \beta \, \text{Unrest}_{il} + \gamma \, \text{Post}_t + \delta \, \text{Unrest}_{il} \times \text{Post}_t + \epsilon_{ilt} \tag{12}$$

where Unrest_{il} is an indicator variable equal to 1 if individual i in area l dealt with an episode of social unrest and Post_{lt} is an indicator variable equal to 1 during and after the episode. However, this specification may be a poor fit for this context. This approach is meant to pick up a policy change, where the policy switches "on" in perpetuity (i.e., treatment is an absorbing state). However, here we see a brief uptick in social unrest, a change which switches "on" and then "off" again. Therefore, in the case where we are estimating the instantaneous treatment effect, we allow treatment to switch off. We can re-write the specification as

$$y_{ilt} = \alpha + \beta \, \text{Unrest}_{il} + \gamma \, \text{Episode}_{lt} + \delta \, \text{Unrest}_{il} \times \text{Episode}_{lt} + \epsilon_{ilt} \tag{13}$$

where Episode_{lt} is an indicator variable (equal to 1 when unrest is taking place). We will refer to this as the DiD specification. Finally, we might estimate this specification with two-way fixed effects, which we will refer to as the TWFE specification:

$$y_{ilt} = \beta_i + \gamma_t + \delta \, \text{Unrest}_{il} \times \text{Episode}_{lt} + \epsilon_{ilt}. \tag{14}$$

However, a non-absorbing treatment is not the only wrinkle we face in estimating our parameter of interest. Our context features variation in the timing of social unrest, which has implications for the estimation of treatment effects when they are heterogeneous across units of time (Goodman-Bacon, 2019; de Chaisemartin and D'Haultfœuille, 2020). The treatment effect estimated using this TWFE specification will be a weighted average of treatment effects computed from the various treatment and comparison groups. In the context of our event-only DiD, and TWFE estimator, this would create comparison groups where people faced with unrest at a later date would be compared to those who face *ongoing* unrest.[23] Using a treatment group for comparison for a newly treated group will produce an underestimate of the treatment effect for that comparison. Moreover, the weights themselves are often unreasonable and an artifact of the estimation procedure. First, weights on these cell specific treatments are sometimes negative, which could result in a treatment effect that is opposite in sign from a cell specific effect. Second, groups treated mid-panel receive higher treatment weights despite the lack of any theoretical reason why this would be the case.

In addition to demonstrating the issue of negative wights, de Chaisemartin and D'Haultfœuille (2020) demonstrates a related issue when heterogeneous treatment effects are at play. In particular, even when all weights are positive, if there is significant heterogeneity in the treatment effects of different groups, it is possible that the average of these group level treatment effects could be opposite the sign of the TWFE estimate. This brings our attention back to the weights generated by TWFE because they are not informed by theory, but rather by an artifact of the estimator. The authors provide two tests that can be used to diagnose the potential for these issues (which we will describe in Section 5.1.1) and a new estimator, DID^M, appropriate for estimating the ATE in our context. This approach is robust to issues brought on by variable treatment timing in combination with heterogeneous and/or dynamic treatment effects. Finally, their approach is a rare estimator in the family of difference-in-differences that not only allows for treatment to turn "on" and "off," but naturally estimates the effect of entering and leaving treatment to construct the treatment effect estimate. However, these features are bought with a number of assumptions, which we will describe in the next subsection.

4.2 DID^M Estimator

The DID^M estimator is the weighted sum of the group level differences that arise when groups move in or out of treatment. A number of conditions and assumptions need to hold for DID^M to deliver valid estimates of the treatment parameter. Some of these conditions involve restrictions on the data structure and can be quickly (and directly) verified. Others must be maintained.

There are three assumptions which we can quickly verify. First, the dataset must feature a balanced panel of groups, as ours does. Second, treatment must be sharp within these groups. That is, all individuals in a group must have the same treatment status. We construct groups based on their most used tower in order to ensure groups are non-overlapping. Third, the data features stable groups to serve as comparisons

---

[23]An instructive example is presented in Goodman-Bacon (2019): in a simple "early and late" treatment model (i.e., group 1 is treated earlier, group 2 later, group 3 not at all), this can be expressed as the weighted sum of four treatment effects (or comparisons). These are: early v. untreated, late v. untreated, early v. late, before treatment and late v. early group after treatment. This last effect is important: groups who were already treated act as controls even after they have received treatment.

for those who switch treatment statuses. That is, when a tower enters treatment there exists another that remains untreated in both periods, and vice versa. We check this in the data and find that there are always untreated groups to compare to when a tower moves into treatment and there xists groups which remain treated in some periods where social unrest lapses. Where there are not stable treatment groups to estimate the movement of an area out of treatment, this data does not contribute to the estimator.

However, to estimate DID^M we must maintain another five assumptions that cannot be directly verified. The first two assumptions involve strong exogeneity. In particular, for the contribution to the ATE of those who join treatment to be identified, strong exogeneity must hold for control towers.[24] For example, this condition forbids cases where protest or other social unrest occurs because of another shock specific to this location at this time. Likewise, the contribution to the ATE of those who leave treatment depends on the same condition for treated towers – social unrest cannot lapse because of another shock.[25]

The next two assumptions involve the parallel trends assumption, which remains crucial as it does in standard DiD. Again we consider symmetric assumptions for the identification of contributions of those who join and those who leave treatment. For those who join treatment, the standard common trends assumption applies.[26] We document these pre-treatment trends in section 5.3.2. Likewise, for those who leave treatment, we invert this common trends assumption to consider common trends among treated areas.[27] Because of our short windows of treatment, however, it is difficult to provide suggestive evidence around the plausibility of this assumption.

Finally, our last assumption to establish causality is a mean independence assumption. More specifically, we assume mean independence between a group's outcome and other groups' treatments. This assumption serves to exclude spillovers from social unrest.[28]

In addition to technical assumptions needed to establish causality, interpreting the treatment effect is similarly important. A key fact in this interpretation is that the protests were disruptive to society but do not directly impact the infrastructure of the cell phone network, so the degree of damage to towers was minimal. This closes off any supply side channel that might impact usage and isolates treatment to the demand side.

4.3 Data Processing, Estimation, and Standard Errors

After cleaning, we process transactions from voice calls made in Haiti in January and February 2019 into a set of network graphs and a list of users.

---

[24]In technical terms, this means that where l is group and t is time, for all $(l,t) \in \{1,\ldots,L\} \times \{2,\ldots,T\}, E(Y_{l,t}(0) - Y_{l,t-1}(0)|D_{l,1},\ldots,D_{l,T}) = E(Y_{l,t}(0) - Y_{g,t-1}(0))$.

[25]Similarly, we can write for all $(l,t) \in \{1,\ldots,L\} \times \{2,\ldots,T\}, E(Y_{l,t}(1) - Y_{l,t-1}(1)|D_{l,1},\ldots,D_{l,T}) = E(Y_{l,t}(1) - Y_{g,t-1}(1))$.

[26]That is, for $t \geq 2, E(Y_{l,t}(0) - Y_{l,t-1}(0))$ does not vary across groups l and particularly across groups who are treated and are not. Note that this assumption does not admit anticipation effects, a fact that drives us to omit protest from our measure of treatment with social unrest.

[27]That is, we assume that for $t \geq 2, E(Y_{l,t}(1) - Y_{l,t-1}(1))$ does not vary across l.

[28]Formally, for all $l, t, E(Y_{l,t}(0)|\mathbf{D}) = E(Y_{l,t}(0)|\mathbf{D}_g)$ and $E(Y_{l,t}(1)|\mathbf{D}) = E(Y_{l,t}(1)|\mathbf{D}_g)$. While we cannot test this assumption, we could document suggestive evidence by analyzing the sensitivity of treatment effects to the distance at which a tower becomes treated. For the moment, we leave this as an future exercise.

4.3.1 Data Cleaning

To clean the data, we remove all transactions involving irregular length of prefix (i.e., length 3 as opposed to 5). Second, we drop all transactions that do not start with a standard 509 Haiti calling code. This has the sum effect of removing all international calls as well as those related to information, emergency, and corporate numbers. Finally, we remove any number that is not associated with a cellphone tower, which are likely landlines, but cannot be observed as calling in the dataset.

4.3.2 Assignment of Nodes to Treatment

We focus on social unrest taking place in the greater Port-Au-Prince region. Particularly in February, social unrest was centered in Haiti's largest city and was geographically differentiated within the city. Additionally, areas outside of Port-Au-Prince do not make good comparisons to the city. Finally, our data quality both in terms of cell service coverage and social unrest events likely do not do justice to the areas outside of Port-Au Prince. Using this logic, we filter cellphone towers to be within the greater Port-Au-Prince region.[29] Next, we match these towers to protest events that happened within one kilometer of the tower. We view this as a relatively inclusive band, but a reasonable one. In particular, it roughly captures the differences between thoroughfares in Port-Au-Prince. Since events of social unrest tend to occur on these thoroughfares, this should capture only one thoroughfare at a time within its surrounding neighborhood. Additionally, since not every event is perfectly precise, we feel that this distance does a good job including those in the actual radius of the event. Using this data, we create a list of treatments by tower and by date, which we merge to the user-tower file. We aggregate the user-tower dataset over the three week baseline period, recording all towers used by a given user. Then, we match this user-tower file to the tower list featuring treatment status. Finally, we assign treatment status to an individual if an unrest event happened at their most used tower. To record a proxy for mobility, we compute both the HHI of call distribution among towers and the share of calls at the most used tower.[30]

4.3.3 Constructing Nodes, Edgelists and Network Statistics

Starting with the clean transaction data, for each day in February, we construct two summary files. First, we construct user-tower files, aggregating the number of calls a user made from a given tower or received at the same tower over the course of that day. Users will serve as nodes in the network, but the tower is preserved to assign them to "treatment" later. In this file, we record the number of calls made and received each day. Second, we construct edgelists—records of all calls between two numbers over the course of that day. In these edgelists, we compute three undirected measures of the network: if the two users called each other, how many times they called one another, and the total duration of these calls. Finally, we compute network statistics for each node using these three measures. In particular, for each node and each day,

---

[29]To do this, we select sub-commune boundaries, also known as communal sections (ADM4) to capture the urban core of Port-au-Prince and its direct surroundings. These sections include Petit Bois, Varreux (Croix-des Bouquets), Bizoton, Rivière-Froide, Thor, Bellevue, Saint Martin, Varreux (Delmas), Bellevue Chardonnieres, Etang du Jong, Martissant, Morne L'Hopital, and Turgeau. The selected communal sections can be seen in figure 3, for example.

[30]Notably, past work with call detail records in Port-au-Prince suggests mobility is low. In particular, Zagatti et al. (2018) find that only 42% of people in the Port-au-Prince area travel more than 1km to work.
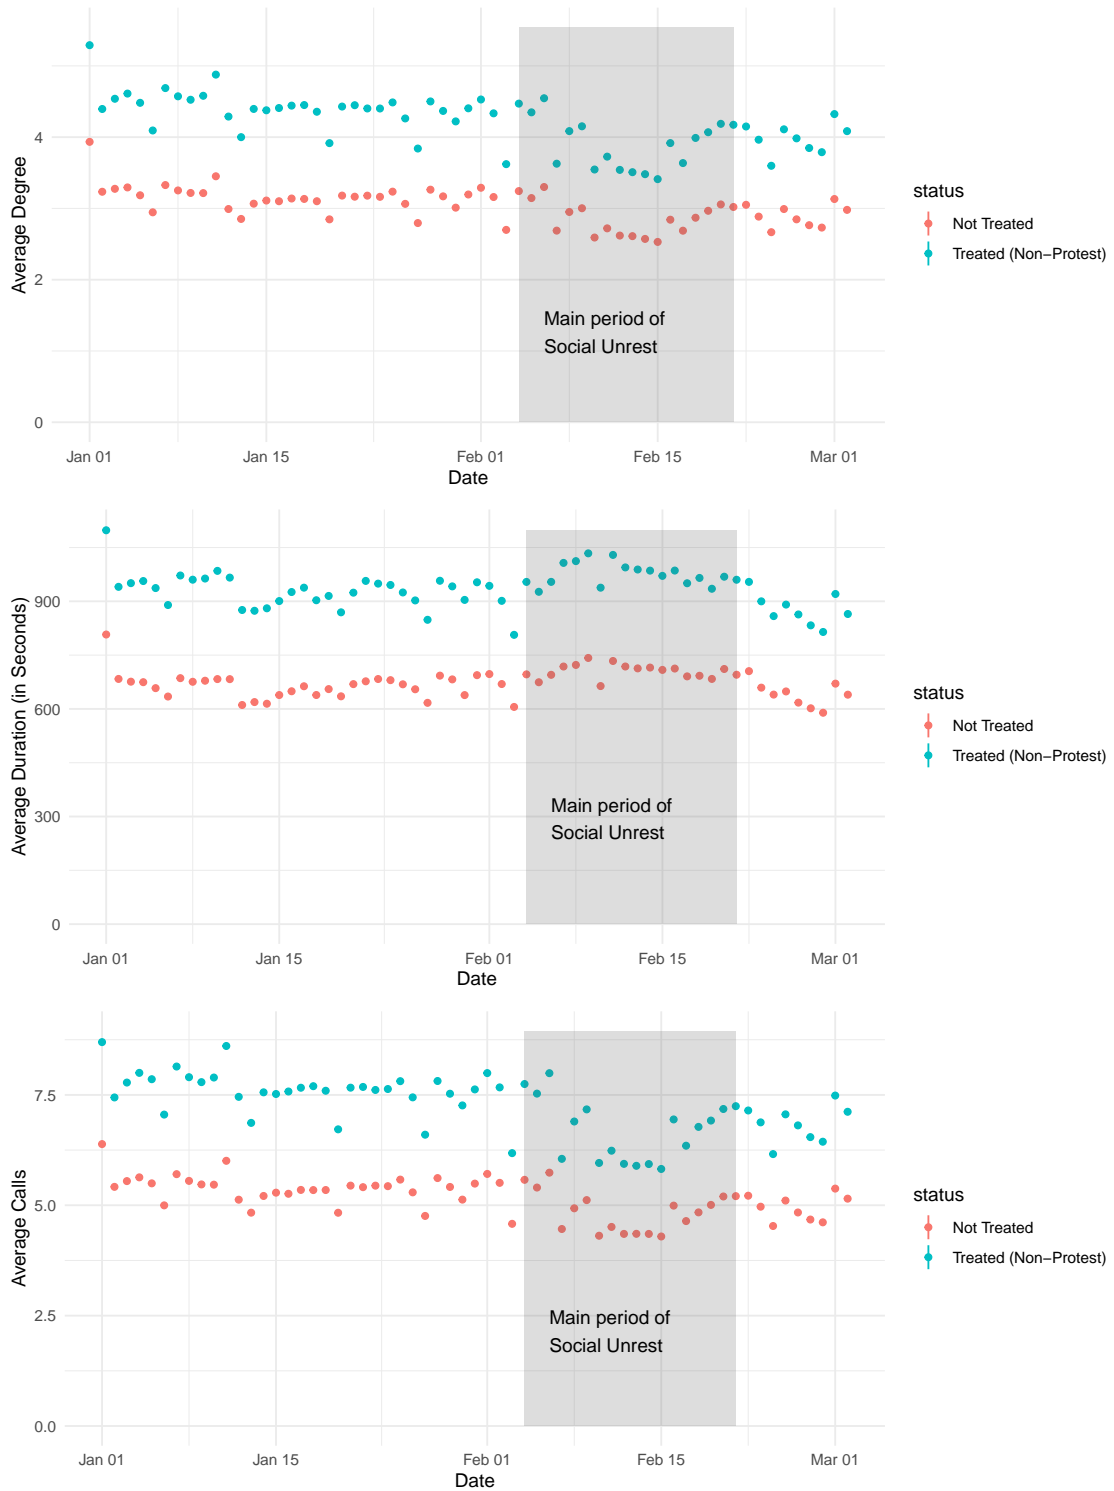
Figure 5: Average degree, duration, and calls over the sample period by exposure to non-protest social unrest event.

we compute network degree by summing up usage by that node—first unweighted, and then weighted by number of calls and by call duration. Taking a list of all users who make calls within the Port-Au-Prince area, we merge treatment and network statistics to these users day by day and compile these into a panel. Because networks are constructed at a country-wide level, this does not drop calls or texts to those outside of Port-Au-Prince.

4.3.4 Estimation and Standard Errors

All estimates are obtained in R. All fixed effects specifications are estimated in R using the command felm from the lfe package. For fixed effects specifications, standard errors are clustered at the tower level (multiple antennas in the same location are considered one tower). The choice of clustering variable follows from the design based approach in Abadie et al. (2017), since treatment is assigned by location of the most used tower. DID^M is estimated using the command did_multiplegt from the package DIDmultiplegt. For DID^M, standard errors are computed using a block bootstrap with tower as the clustering variable (500 repetitions). Finally, diagnostic tests of the TWFE weights are performed using the command twowayfeweights from the package TwoWayFEWeights.

5 Results

5.1 Main Effects: Non-Protest Social Unrest Events

5.1.1 Diagnostic Tests of the TWFE Estimator

To diagnose robustness of TWFE to heterogeneous treatment effects, we run two tests outlined by de Chaisemartin and D'Haultfœuille (2020) in Corollary 1. We run these tests under a common trends assumption, using our non-protest social unrest variable as treatment and degree, call weighted degree, and duration weighted degree as outcomes. Because of the computational expense of the DID^M estimator, we run these tests on a subsample of 100,000 members of the population, and over five weeks, beginning the 21st of January until the 24th of February. For all three outcomes, we find similar results when characterizing the weights of TWFE. All ATTs estimated by TWFE are assigned positive weights, and therefore, the sum of the positive weights is equal to 1. This result eliminates one issue that we commonly test for, which is that the TWFE coefficient could plausibly be of a different sign than all of the ATTs. However, the estimate produced by is still compatible with a data generating process where the average of those ATT is equal to 0, meaning changes in weights could lead to reversals in the sign of the treatment effect. This suggests there is reason to be concerned about the TWFE weighting, given the fact that the TWFE estimator yields weights that may be arbitrary.

5.1.2 Comparison of Estimators

We estimate treatment effects using the DiD, TWFE, and DID^M on a common subsample of 100000 users over five weeks, beginning the 20th of January until the 24th of February. Using both the TWFE estimator and the DID^M estimators, we see that while the number of contacts and calls per day falls, duration spent

Table 3: Effects of Social Unrest on Network Degree

| | Degree (Total Contacts) | | |
	DiD	TWFE	DID^M
Episode of Non-Protest Social Unrest	0.021	-0.085^{***}	-0.048^{***}
	(0.031)	(0.025)	(0.012)
Ever Non-Protest Social Unrest	0.259**		
	(0.102)		
Day FE	Yes	Yes	
Indv. FE	No	Yes	
Cluster	Tower	Tower	Tower†
$N \times T$	3,067,920	3,067,920	3,067,920
R^2	0.005	0.648	
Adjusted R^2	0.005	0.638	
Residual Std. Error	5.094	3.071	
df RSE	3,067,882	2,982,664	
N effect			1,065,149
N switchers effect			181,124

Notes:

***Significant at the 1 percent level.
**Significant at the 5 percent level.
*Significant at the 10 percent level.
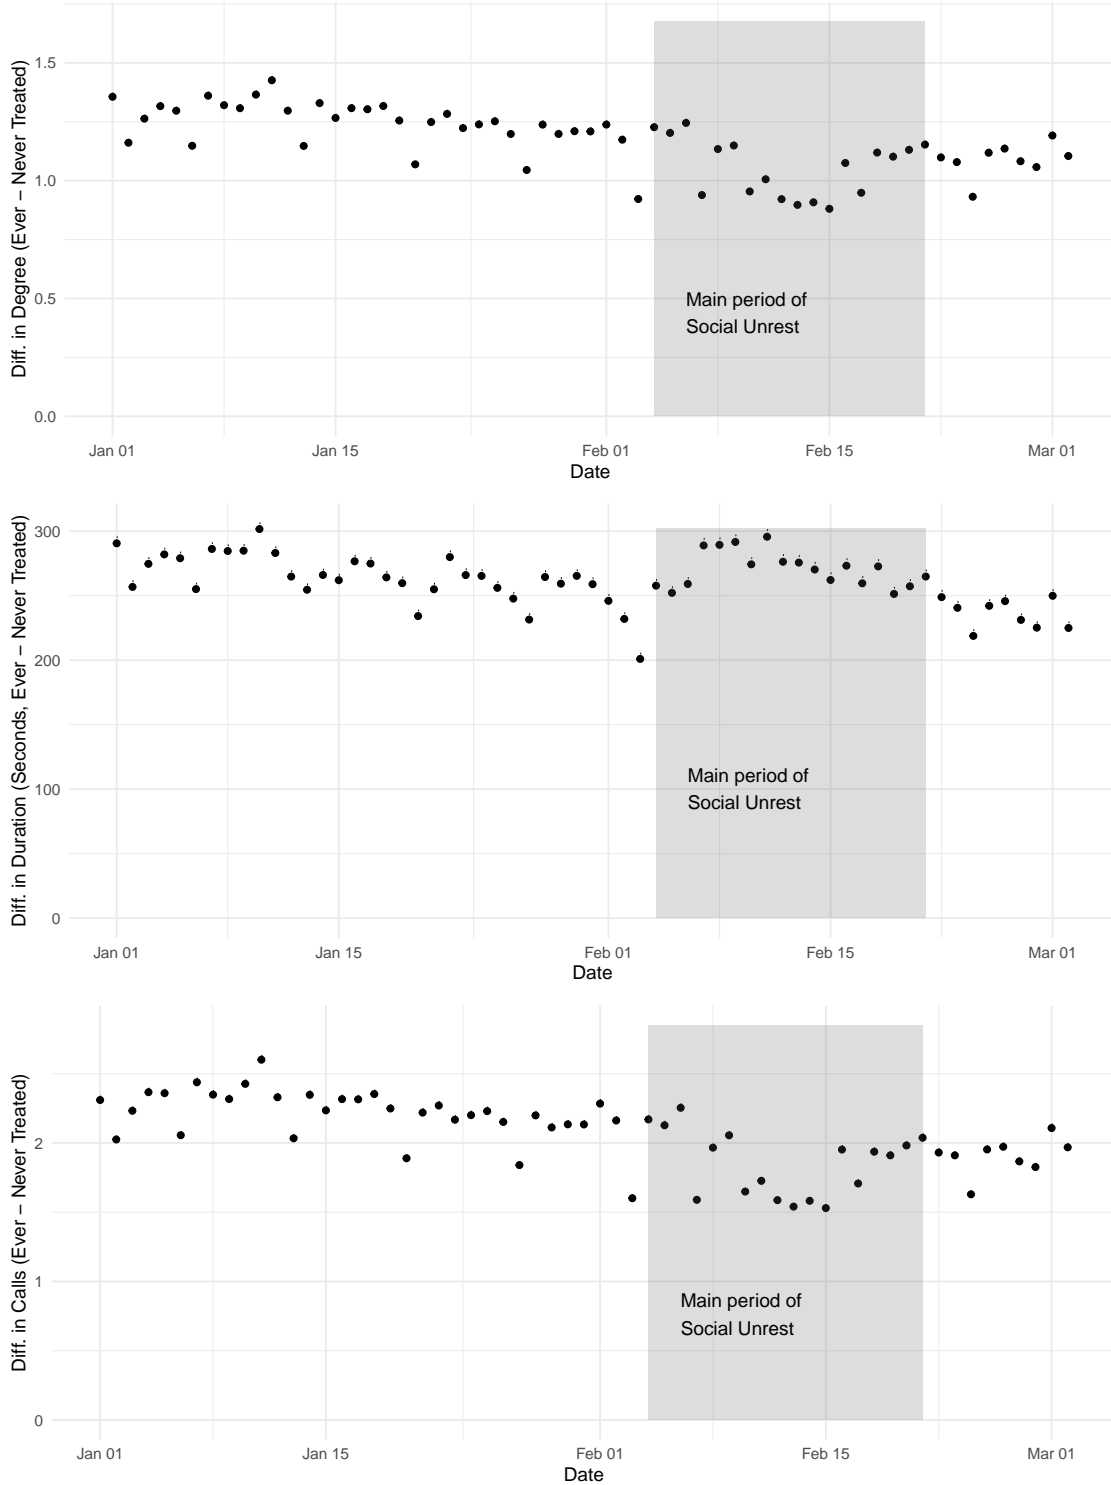† Block bootstrap with 500 repetitions

Figure 6: Differences in degree, duration, and calls over the sample period by exposure to non-protest social unrest event.

conversing does not. In particular, using duration is positively but insignificantly impacted whereas using DID^M duration is negatively impacted but still insignificant. Therefore, based on these results, we conclude that total duration remains similar to before.

Despite the concordance in sign between the DID^M estimates and the TWFE estimates on contacts and calls within this sample, our estimates differ in magnitude across these estimation methods. While coefficients are on the same order of magnitude, the differences in results do matter in a quantitative sense. Taking the DID^M estimates as the true effect, we see bias in the TWFE estimates for contacts and calls on the order of 48 and 77% of the size of the DID^M coefficients. Likewise, the coefficient on duration switches signs when moving from TWFE to DID^M. We interpret these differences in coefficient as reflecting differences in weighting between the TWFE and DID^M estimators. Given the known problem with TWFE weights when treatment effects are heterogeneous and there is variation in treatment timing, we should prefer the weighting for the DID^M estimator which is intentional as opposed to arbitrary. Therefore, given these differences, we select DID^M as our preferred estimator.[31]

5.1.3 Estimates and Interpretation of Causal Effects

Using DID^M we estimate episodes of social unrest reduce degree by 0.048 contacts, but do not cause people to spend significantly less time talking on the phone. Considered another way, individuals spend *more* time talking per contact in periods of social unrest: the reduction in contacts amounts to roughly one percent of average daily contacts over the period, while the (insignificant) reduction in duration is about one quarter of one percent.

Notably, the reduction in number of contacts is consistent with our hypothesis from the model that people cease communication with those with whom they have weak ties. While the formal model does not take a stance on node level duration, our hypotheses based on assumptions about the size of informational shocks, suggested we should see an increase in duration for those treated with social unrest. These results are therefore inconsistent with our hypothesis but not the formal model overall.

Interestingly, this pattern of result bears resemblance to results in related contexts. Romero et al. (2016) considers the metadata of text messages sent in a hedge fund during price shocks, using a form of TWFE to estimate that network structure becomes more tightly knit in the face of such shocks.[32] Likewise, Blumenstock et al. (2016), finds that after a large earthquake in Rwanda, calls and airtime transfers are made to those in the affected area. In particular, transfers are made between pairs of individuals with histories of reciprocal favor exchange. Finally, Jia et al. (2021) finds that after an earthquake in the Yu'an province, calls are more likely to be made to those in their family when families were closer knit. Like in those settings, this pattern of results might imply people interacting more with their close or reciprocal contacts.

However, the precise zero on duration also raises another possibility of a nuance that departs modestly

---

[31]When errors are homoskedastic and uncorrelated, TWFE may yield a lower variance estimate than DID^M even though it is biased. Thus, there would be a bias-variance trade-off between TWFE and DID^M. While our errors (and treatment) are correlated, this bias-variance trade-off does not play out in our case. In particular, DID^M sees smaller standard errors in some cases.

[32]While users are facing stress as well, our setting clearly differs in that most cellphone users will not feel culpable for previously made decisions when social unrest strikes. Hedge fund employees on the other hand may need to wrestle with assignments of blame if losses are accrued during such price shocks.

Table 4: Effects of Social Unrest on Total Call Duration

	Total Duration (in seconds)		
	DiD	TWFE	DID^M
Episode of Non-Protest Social Unrest	36.813***	9.974	-4.448
	(9.088)	(6.282)	(6.990)
Ever Non-Protest Social Unrest	87.738***		
	(21.979)		
Day FE	Yes	Yes	
Indv. FE	No	Yes	
Cluster	Tower	Tower	Tower†
$N \times T$	3,067,920	3,067,920	3,067,920
R^2	0.001	0.507	
Adjusted R^2	0.001	0.493	
Residual Std. Error	1,819.503	1,295.769	
df RSE	3,067,882	2982664	
N effect			1,065,149
N switchers effect			181,124

Notes: ***Significant at the 1 percent level.
**Significant at the 5 percent level.
*Significant at the 10 percent level.
† Block bootstrap with 500 repetitions

from the theoretical model we have presented earlier. In particular, if budget constraints bind or it is not possible to obtain airtime credit or visiting an airtime vendor, it may be that there is a process of substitution taking place around events of social unrest. That is, while similar to the story above, the move from very weak ties to stronger or more informed ties is navigated via a budget constraint.

5.2 Decomposing Results by Tie Strength and Information

While these results paint a picture consistent with our theoretical model, more extensive testing of the model hypotheses is needed. In this section, we provide evidence around strong and weak ties as well as high and low information ties. We decompose the results by restricting what relationships are included when we construct outcomes. For example, strong and informed tie degree would computed by adding only the contacts who were both strong ties and had high-degree at baseline.[33] If we find more positive (or less negative) effects of social unrest on these restricted outcomes, this would provide evidence that people people value these connections in uncertain situations. These estimates decompose the treatment effect into peer subgroups. By construction, the sum of the subgroup effects equal the effect of all subgroups.

---

[33]To recall, strong ties are those alters with dyadic duration above 80th percentile in the three week baseline period. Informed alters are those have degree above 67th percentile in the baseline period.

Table 5: Decomposing DID^M Estimates of Social Unrest on Degree

Restrictions	DID^M Estimates		Pre-Unrest[‡]		Scaled Effect[§]
	Effect	Std Err.[†]	Mean	Std Dev.	
(None: Main Effect)	-0.0478	0.0145	3.107	4.418	-0.015
Strong Ties (Any Info)	-0.0340	0.0090	1.274	2.224	-0.027
Informed Ties (Any Strength)	-0.0201	0.0067	1.833	2.819	-0.011
Strong & Informed Ties	-0.0221	0.0060	0.798	2.145	-0.028
Strong & Uninformed Ties	0.0020	0.0030	0.475	1.052	0.004
Weak & Informed Ties	-0.0118	0.0050	0.563	1.602	-0.021
Weak & Uninformed Ties	-0.0158	0.0091	1.270	1.951	-0.012

Notes:

[***]Significant at the 1 percent level.

[**]Significant at the 5 percent level.

[*]Significant at the 10 percent level.

[†] Block bootstrap with 500 repetitions

[‡]First ten days of data, pre-main period of social unrest

[§] Scaled Effect = Effect / Pre-Unrest Mean

For example, the sum of effects on strong ties and weak ties will equal the main (unrestricted) effect.

Table 5 presents results for degree (total contacts). While total contacts fall, this is not evenly distributed across subgroups. In particular, for those connections who are strong and uninformed ties, we actually find that contacts do not fall. This suggests that people want information about their close friends, family, or associates. We also find a smaller reduction compared to baseline for people's weak uninformed ties, an effect that may require greater exploration. Finally, informed ties (both strong and weak) have a relatively large reduction in terms of contacts. This may provide evidence against the "strength of weak ties" in this case (Granovetter, 1973). Table 6 presents the decomposition of effects on duration of calls. As is the case with the main effect, the subgroup effects are null results. This pattern of results reflects the search for a very limited form of information. Instead of learning about the state of the world generally, individuals seek information about realities for those people that are close to them and recede from searching for other information.

5.3 Addressing Threats to Validity

5.3.1 Network Response and External Validity

A common concern when studying social network hypotheses using digital trace data is the external validity of estimated responses (Howison et al., 2011). These include issues related to network sampling and the relationship between in person networks and ICT-mediated social networks.

A first potential threat to validity can be related to network sampling. Namely, the CDR data does not cover the universe of households in Port-Au-Prince. However, because we have administrative data from

Table 6: Decomposing DID^M Estimates of Social Unrest on Total Call Duration

Restrictions	DID^M Estimates		Pre-Unrest[‡]		Scaled Effect[§]
	Effect	Std Err.[†]	Mean	Std Dev.	
(None: Main Effect)	-4.45	5.86	663.6	1484.1	-0.007
Strong Ties (Any Info)	0.96	5.08	436.2	1266.6	0.002
Informed Ties (Any Strength)	0.56	3.91	297.4	1063.3	0.002
Strong & Informed Ties	2.07	3.53	241.6	978.2	0.009
Strong & Uninformed Ties	-1.11	3.13	194.6	861.1	-0.006
Weak & Informed Ties	-1.51	1.39	55.7	228.5	-0.027
Weak & Uninformed Ties	-3.90	3.55	171.6	515.2	-0.023

Notes:

[***]Significant at the 1 percent level.

[**]Significant at the 5 percent level.

[*]Significant at the 10 percent level.

[†] Block bootstrap with 500 repetitions

[‡]First ten days of data, pre-main period of social unrest

[§] Scaled Effect = Effect / Pre-Unrest Mean

a telecom provider with a large market share, this should allow for census-like conditions (Chandrasekhar and Lewis, 2016) . Therefore, we argue we can interpret network effects as though nodes are not sampled.

Another potential issue is that of how ICT-mediated social networks relate to in-person social networks. Notably, past work studying the relationship between egocentric network structure found in online and in-person networks has shown similarities across a broad range of structures and relationship types (Dunbar et al., 2015; Bisbee and Larson, 2017). However, response in ICT-mediated social networks might also differ from that of in-person networks (for example, due to substitution between the two networks). If so, we may worry that estimated effects are biased. While we cannot rule this out, it is important to note that to estimate effects which are qualitatively wrong (i.e., which are of opposite sign) communication in ICT-mediated social networks and in-person would need to both move in opposite directions, and there would need to be a larger effect in in-person communication. It is not clear why these would move in opposite directions *per se*. While unrest may lead to reduced mobility, in Port-Au-Prince's dense urban core this should not reduce ability to communicate in-person on the day of unrest.

In interpreting our results as saying something about social networks, we rely on the fact that communications networks are a good proxy for the true social networks of interest that may exist over multiple "platforms." One issue that might arise in particular is that the use of smartphones has grown in Haiti over the past few years. Therefore, we may not observe the full effect of social unrest on digital communication due to app based calling and messaging apps (e.g., WhatsApp). We might be able to observe a proxy for smartphone messaging over time by looking at cellphone data usage, but we are not able to observe communication that take place via WhatsApp and other messaging services. Despite this, we can at least provide evidence around these issues by building a restricted subsample of only those who tend to use

largely voice or SMS services and less data than one would need to utilize a smartphone.

5.3.2 Pre-Treatment Trends

To provide qualitative evidence about the plausibility of the pre-trends assumption, we visualize our data over the two month study period. Figure 5 plots the averages of outcomes by exposure to non-protest social unrest events. To demarcate the most intense period of social unrest, we shade the area after which the number of social unrest incidents spiked in Port-Au-Prince.

A number of patterns within the pre-trends are interesting. Usage is higher among those in areas that were ever treated, suggesting that these social unrest events were associated with some feature of the area. It is difficult to say exactly what drives this, though proximity to roads or gathering spaces, which tend to be used for commerce. For the three outcomes, we see reductions in calling on Saturday and Sunday. This is most pronounced when considering degree or calls, and less pronounced when considering duration. However, calling behavior follows remarkably similar trends between areas where non-protest social unrest took place and those where it did not.

Similar figures are included in the appendix for two other definitions of treatment. First, figure 14 presents these same outcomes by exposure to protest. Second, figure 15 presents the outcomes by exposure to any type of social unrest (protest or non-protest). Qualitative conclusions from inspection of these figures is broadly similar. Notably, however, including all social unrest events in treatment magnifies the difference int pre-trend outcomes as compared to the other two definitions.

While formal tests of pre-trends are often undertaken, such tests can further bias conventional estimates when the pre-trends assumption is violated (Roth, 2022). Therefore, we opt not to report pre-trends tests until we are able to account for the distortions from pre-testing in our estimates.

6   Conclusion

In this paper, we estimate the treatment effect of social unrest on social networks using call networks in Haiti as a proxy for social networks. While significant work is needed to refine the estimation and interpretation of these effects, some interesting preliminary results appear. Consistent with a model of calling during social unrest, we find that network degree falls along with the number of calls made while the duration of time spent talking remains constant. These results paint a story where people reach out their close friends (or their close friends reach out to them) when their neighborhood is the sight of social unrest. Decomposing these results, we find evidence the people check in on these close ties but do not reach out to individuals who, due to their centrality at baseline, might be more informed about the crisis at hand.

Just as more evidence around what types of connections are valued is needed, assessing other threats to the validity of our estimates is a pressing concern going forward. First, while we have a large and random sample of the population of interest, testing the same hypotheses in additional subsamples might result in different in results due to sampling error. Ideally, we would not need to sample at all, though the DID^M is very memory intensive, thus requiring this measure. However, we are able to use substantially larger

samples when only computing the TWFE estimator, potentially even a full sample. Therefore this could serve as a measure for how much these individual samples differ from the population sample.

Second, to further assess pre-treatment trends, de Chaisemartin and D'Haultfœuille (2020) presents a placebo estimator which constructs placebo treatment effects from the time periods just before individuals switch into treatment, therefore testing for pre-trends. Additionally, Roth (2022) presents methods of testing pre-treatment trends and then accounting for this within the main specification.

Third, we may want to account for heterogeneity in treatment effects via mobility of cellphone users. Using baseline estimates of tower diversity and call share at their primary tower, we can characterize the importance of mobility on the results at hand. Using these results, we may be able to address some attenuation that results from our method of assigning treatment.

Fourth, while we define treatment as non-protest social unrest episodes, it would be interesting both to consider descriptive measures of treatment for all social unrest events, and protest only. Given placebo tests discussed earlier and the DID^M estimator, we can then assess how crucial our treatment definition is. While we suppose that protests are anticipated, the placebo estimator could provide evidence for or against this.

Fifth, to better understand treatment assignment, we may also want to vary the distance from tower to social unrest by which the tower becomes treated. In particular, we will construct an alternative definition of treatment where only towers within a smaller radius of social unrest are considered treated. Additionally, this could serve as a source of heterogeneity in treatment. For example, if calling behavior differs between those who are directly next to social unrest as compared to those who are at a safe distance, characterizing these differences might be interesting.

Sixth, all of these outcomes could be similarly computed with alternative measures of the network. In this paper, communication networks, and more specifically call networks, serve as a proxy for social networks. Other measures of communication are available, most notably SMS networks. Future work might investigate these SMS networks using similar approaches. Likewise an analysis of data usage might give clues about the use of alternative calling or messaging services including WhatsApp or Facebook. Similarly the recent outage in WhatsApp services might serve as an interesting setting to study the relevance of these services on network measurement using communications networks.

References

A. Abadie, S. Athey, G. Imbens, and J. Wooldridge. When Should You Adjust Standard Errors for Clustering?, 2017.

D. P. Aldrich and M. A. Meyer. Social Capital and Community Resilience. *American Behavioral Scientist*, 59(2):254–269, 2015. ISSN 15523381. doi: 10.1177/0002764214550299.

A. Ambrus, M. Mobius, and A. Szeidl. Consumption Risk-sharing in Social Networks. *American Economic Review*, 104(1):149–182, 2014. doi: 10.1257/aer.104.1.149.

K. T. Andrews and M. Biggs. The Dynamics of Protest Diffusion: Movement Organizations, Social Networks, and News Media in the 1960 Sit-Ins. 1:752–777, 2006.

Anonymous. Whispers: I Get Paid to Burn Tires and Protest in Haiti. http://woymagazine.com/2015/07/10/paid-burn-tires-protest-haiti/%0A, 2015.

Associated Press. Protesters Stone Haitian President's Home, Battle Police. *VOA*, Feb. 2019.

S. Azid, B. Sharma, K. Raghuwaiya, A. Chand, S. Prasad, and A. Jacquier. SMS based flood monitoring and early warning system. *ARPN Journal of Engineering and Applied Sciences*, 10(15):6387–6391, 2015. ISSN 18196608.

A. Banerjee, E. Breza, A. G. Chandrasekhar, and M. Mobius. Naive Learning with Uniformed Agents. Technical report, 2018.

A. Banerjee, A. G. Chandrasekhar, E. Duflo, and M. O. Jackson. Using Gossips to Spread Information: Theory and Evidence from Two Randomized Controlled Trials. *Review of Economic Studies*, 86(6):2453–2490, 2019a. ISSN 1467937X. doi: 10.1093/restud/rdz008.

A. V. Banerjee, A. G. Chandrasekhar, E. Duflo, and M. O. Jackson. The diffusion of microfinance. *Science*, 341(6144), 2013. ISSN 1095-9203. doi: 10.1126/science.1236498.

A. V. Banerjee, E. Breza, A. G. Chandrasekhar, and B. Golub. When Less is More: Experimental Evidence Based on Information Delivery during India's Demonitization. Technical report, 2019b.

L. A. Beaman, A. BenYishay, J. Magruder, and A. M. Mobarak. Can Network Theory-Based Targeting Increase Technology Adoption? *American Economic Review*, 111(6):1918–1943, 2021. ISSN 0898-2937. doi: 10.2139/ssrn.3225815.

L. Bengtsson, X. Lu, A. Thorson, R. Garfield, and J. von Schreeb. Improved response to disasters and outbreaks by tracking population movements with mobile phone network data: A post-earthquake geospatial study in haiti. *PLoS Medicine*, 8(8):1–9, 2011. ISSN 15491277. doi: 10.1371/journal.pmed.1001083.

D. Bennett, C. F. Chiang, and A. Malani. Learning during a crisis: The SARS epidemic in Taiwan. *Journal of Development Economics*, 112:1–18, 2015. ISSN 03043878. doi: 10.1016/j.jdeveco.2014.09.006.

Q. Bernier and R. Meinzen-Dick. Networks for resilience: The role of social capital. Technical report, International Food Policy Research Institute, Washington, DC, 2020.

J. Bisbee and J. M. Larson. Testing social science network theories with online network data: An evaluation of external validity. *American Political Science Review*, 111(3):502–521, 2017. ISSN 15375943. doi: 10.1017/S0003055417000120.

D. Björkegren. The adoption of network goods: Evidence from the spread of mobile phones in Rwanda. *Review of Economic Studies*, 86(3):1033–1060, 2019. ISSN 1467937X. doi: 10.1093/restud/rdy024.

J. Blumenstock, E. B. Kapstein, and T. L. Scherer. How Do Firms Respond to Insecurity? Evidence from Afghan Corporate Phone Records. 2020.

J. E. Blumenstock, N. Eagle, and M. Fafchamps. Airtime transfers and mobile communications: Evidence in the aftermath of natural disasters. *Journal of Development Economics*, 120:157–181, 2016. doi: 10.1016/j.jdeveco.2016.01.003.

J. E. Blumenstock, G. Chi, and X. Tan. Migration and the value of social networks, 2019.

Y. Bramoullé and G. Genicot. Diffusion Centrality: Foundations and Extensions. Technical Report October, 2018.

L. Bursztyn, D. Cantoni, D. Y. Yang, N. Yuchtman, and Y. J. Zhang. Persistent Political Engagement: Social Interactions and the Dynamics of Protest Movements. *American Economic Review: Insights*, 3(2):233–250, 2021. ISSN 2640-205X. doi: 10.1257/aeri.20200261.

D. E. Campbell. Social Networks and Political Participation. *Annual Review of Political Science*, 16:33–48, 2013. ISSN 10942939. doi: 10.1146/annurev-polisci-033011-201728.

D. Centola and M. Macy. Complex Contagions and the Weakness of Long Ties. *American Journal of Sociology*, 113(3):702–734, Nov. 2007. ISSN 0002-9602, 1537-5390. doi: 10.1086/521848.

D. Centola, V. M. Eguíluz, and M. W. Macy. Cascade dynamics of complex propagation. *Physica A: Statistical Mechanics and its Applications*, 374(1):449–456, Jan. 2007. ISSN 03784371. doi: 10.1016/j.physa.2006.06.018.

A. G. Chandrasekhar and R. Lewis. Econometrics of Sampled Networks, 2016.

J. Charles. Haiti owes Venezuela $2 billion – and much of it was embezzled, Senate report says. *Miami Herald*, Nov. 2017.

J. Charles. For Haitians, a reprieve from violence and protests on Sunday, but uncertainty remains. *Miami Herald*, Feb. 2019.

T. G. Conley and C. R. Udry. Learning about a New Technology: Pineapple in Ghana. *American Economic Review*, 100(1):35–69, 2010. ISSN 0002-8282. doi: 10.1257/aer.100.1.35.

Creole Haiti. 85 haitian creole proverbs that are full of wisdom. https://creole-haiti.com/en/haitian-creole-proverbs/.

C. de Chaisemartin and X. D'Haultfœuille. Two-Way Fixed Effects Estimators with Heterogeneous Treatment Effects. *American Economic Review*, 110(9):2964–2996, 2020. ISSN 19447981. doi: 10.1257/aer.20181169.

J. De Weerdt and S. Dercon. Risk-sharing networks and insurance against illness. *Journal of Development Economics*, 81:337–356, 2006. doi: 10.1016/j.jdeveco.2005.06.009.

V. J. Déralciné and P. Jackson. *Rocks in the Water, Rocks in the Sun: A Memior from the Heart of Haiti.* 2015.

A. Dobra, N. E. Williams, and N. Eagle. Spatiotemporal detection of unusual human population behavior using mobile phone data. *PLoS ONE*, 10(3):1–20, 2015. ISSN 19326203. doi: 10.1371/journal.pone.0120449.

R. I. Dunbar, V. Arnaboldi, M. Conti, and A. Passarella. The structure of online social networks mirrors those in the offline world. *Social Networks*, 43:39–47, 2015. ISSN 03788733. doi: 10.1016/j.socnet.2015.04.005.

N. Eubank and D. Kronick. Friends Don't Let Friends Free Ride. *Quarterly Journal of Political Science*, 16 (4):533–557, 2021. ISSN 1554-0626, 1554-0634. doi: 10.1561/100.00020143.

M. Fafchamps and S. Lund. Risk-sharing networks in rural Philippines. *Journal of Development Economics*, 71:261–287, 2003. doi: 10.1016/S0304-3878(03)00029-4.

Flowminder. Haiti Population movements following the Haiti earthquake on 14 August 2021 and the Tropical Depression Grace , estimated with mobile operator data from Digicel Haiti : Report from 27 August . 1(1):0–6, 2021.

F. Gélineau, J. D. Montalvo, and V. Schweizer-Robinson. Political Culture of Democracy in Haiti and in the Americas 2021: Taking the Pulse of Democracy. Technical report, LAPOP, 2021.

A. Goodman-Bacon. Difference-in-Differences with Variation in Treatment Timing. Technical Report July, 2019.

M. Granovetter. The Strength of Weak Ties. *American Journal of Sociology*, 78(6), 1973. doi: 10.1017/S0269889712000130.

D. Gundogdu, O. D. Incel, A. A. Salah, and B. Lepri. Countrywide arrhythmia: Emergency event detection using mobile phone data. *EPJ Data Science*, 5(1), 2016. ISSN 21931127. doi: 10.1140/epjds/s13688-016-0086-0.

HaitiLibre. Haiti - FLASH : 6th day of paralysis, the country sinks into chaos... - HaitiLibre.com : Haiti news 7/7. Feb. 2019.

J. Howison, A. Wiggins, and K. Crowston. Validity Issues in the Use of Social Network Analysis with

Digital Trace Data. *Journal of the Association for information Systems*, 12(12):767–797, 2011.

W. Jack and T. Suri. Risk Sharing and Transactions Costs: Evidence from Kenya's Mobile Money Revolution. *American Economic Review*, 104(1):183–223, 2014. doi: 10.1257/aer.104.1.183.

M. O. Jackson, T. Rodriguez-Barraquer, and X. Tan. Social capital and social quilts: Network patterns of favor exchange. *American Economic Review*, 102(5):1–45, 2012. doi: 10.1257/aer.102.5.1857.

J. S. Jia, Y. Li, X. Lu, Y. Ning, N. A. Christakis, and J. Jia. Triadic embeddedness structure in family networks predicts mobile communication response to a sudden natural disaster. *Nature Communications*, 12(1), 2021. ISSN 20411723. doi: 10.1038/s41467-021-24606-7.

M. King. The probabilities of node-to-node diffusion in fixed networks. Technical report, 2020.

S. Lemaire. US Deplores Haiti Violence. *VOA*, Feb. 2019.

X. Lu, L. Bengtsson, and P. Holme. Predictability of population displacement after the 2010 Haiti earthquake. *Proceedings of the National Academy of Sciences*, 109(29):11576–11581, 2012. ISSN 00278424. doi: 10.1073/pnas.1203882109.

N. Magnan, D. J. Spielman, T. J. Lybbert, and K. Gulati. Leveling with friends: Social networks and Indian farmers' demand for a technology with heterogeneous benefits. *Journal of Development Economics*, 116: 223–251, 2015. ISSN 03043878. doi: 10.1016/j.jdeveco.2015.05.003.

J.-Y. Merilus. Rural development: The economic potentials of haiti's "lakou" system. *Focus on Geography*, 58(1):36–45, 2015. doi: 10.1111/foge.12047.

D. Metaxa-Kakavouli, P. Maas, and D. P. Aldrich. How Social Ties Influence Hurricane Evacuation Behavior. *Proceedings of the ACM on Human-Computer Interaction*, 2(CSCW):1–16, Nov. 2018. ISSN 2573-0142. doi: 10.1145/3274391.

J. Miller. *Redefining the Lakou: The Resilience of a Vernacular Settlement Pattern in Post-Disaster Haiti.* PhD thesis, University of Oregon, 2013.

M. Mooijman, J. Hoover, Y. Lin, H. Ji, and M. Dehghani. Moralization in social networks and the emergence of violence during protests. *Nature Human Behaviour*, 2(6):389–396, 2018. ISSN 23973374. doi: 10.1038/s41562-018-0353-0.

O. Oh, K. H. Kwon, and H. R. Rao. An Exploration of Social Media in Extreme Events: Rumor Theory and Twitter During the Haiti Earthquake 2010. In *ICIS 2010 Proceedings*, volume 231, 2010.

R. B. Patel and K. M. Gleason. The association between social cohesion and community resilience in two urban slums of Port au Prince, Haiti. *International Journal of Disaster Risk Reduction*, 27:161–167, Mar. 2018. ISSN 22124209. doi: 10.1016/j.ijdrr.2017.10.003.

A. V. Pierre-Charles, T. J. Lybbert, D. S. Putman, and O. B. Cabanillas. 2019 Haiti Social Unrest Timeline, 2020.

G. J. Rahill, N. E. Ganapati, J. C. Clérismé, and A. Mukherji. Shelter recovery in urban Haiti after the earthquake: The dual role of social capital. *Disasters*, 38(s1):S73–S93, Apr. 2014. ISSN 03613666. doi: 10.1111/disa.12051.

E. Riley. Mobile money and risk sharing against village shocks. *Journal of Development Economics*, 135 (June):43–58, 2018. ISSN 03043878. doi: 10.1016/j.jdeveco.2018.06.015.

D. M. Romero, B. Uzzi, and J. Kleinberg. Social Networks Under Stress. In *WWW '16: Proceedings of the 25th International Conference on World Wide Web*, 2016. ISBN 978-1-4503-4143-1. doi: 10.1145/2872427.2883063l.

J. Roth. Pretest with Caution: Event-Study Estimates after Testing for Parallel Trends. *American Economic Review: Insights*, 4(3):305–322, Sept. 2022. ISSN 2640-205X, 2640-2068. doi: 10.1257/aeri.20210236.

T. Sakaki, M. Okazaki, and Y. Matsuo. Earthquake Shakes Twitter Users: Real-time Event Detection by Social Sensors. In *WWW '10*, pages 851–860, 2010. ISBN 978-1-60558-799-8.

L. See. A review of citizen science and crowdsourcing in applications of pluvial flooding. *Frontiers in Earth Science*, 7(March):1–7, 2019. ISSN 22966463. doi: 10.3389/feart.2019.00044.

E. J. Sommerfeldt. Disasters and Information Source Repertoires: Information Seeking and Information

Sufficiency in Postearthquake Haiti. *Journal of Applied Communication Research*, 43(1):1–22, Jan. 2015. ISSN 0090-9882, 1479-5752. doi: 10.1080/00909882.2014.982682.

Z. C. Steinert-Threlkeld, D. Mocanu, A. Vespignani, and J. Fowler. Online social networks and offline protest. *EPJ Data Science*, 4(1):1–9, 2015. ISSN 21931127. doi: 10.1140/epjds/s13688-015-0056-y.

S. H. Strogatz and D. J. Watts. Collective Dynamics of 'Small-World' Networks. *Nature*, 393(June):440–442, 1998.

G. A. Zagatti, M. Gonzalez, P. Avner, N. Lozano-Gracia, C. J. Brooks, M. Albert, J. Gray, S. E. Antos, P. Burci, E. zu Erbach-Schoenberg, A. J. Tatem, E. Wetter, and L. Bengtsson. A trip to work: Estimation of origin and destination of commuting patterns in the main metropolitan regions of Haiti using CDR. *Development Engineering*, 3(October 2017):133–165, 2018. ISSN 23527285. doi: 10.1016/j.deveng.2018.03.002.

A Theoretical Appendix

A.1 Useful Properties of the Utility Function

We build on the theoretical model from Björkegren (2019). In particular, that paper uses the utility function

$$v(d, \varepsilon) = d - \frac{1}{\epsilon} \left[\frac{d^\gamma}{\gamma} + \alpha \right]. \tag{15}$$

Our adaption of this utility function features six useful properties:[34]

1. Zero call duration yields zero utility, $v(0, \epsilon) = 0$

2. Diminishing marginal returns to call duration, $v(d, \epsilon)$, is concave in d

3. For some values of the parameters a call is placed; The optimal duration yields non-negative utility, $v(d^*, \epsilon) \geq 0$ where d^* solves $\frac{\partial v(d^*, \epsilon)}{\partial d} = c$ or is zero.

4. Even if calls are free, a caller won't talk forever. That is, even when marginal cost equals zero, there is a duration d^* where $\frac{\partial v(d^*, \epsilon)}{\partial d} = 0$

5. Changing the cost of a call changes the extensive decision to call. This requires the marginal utility of calling to be finite at zero. $\frac{\partial v(0, \epsilon)}{\partial d} < \infty$.

6. Changing the marginal cost of a call affects longer calls more than shorter calls, $\frac{\partial^2 d^*}{\partial c \partial \epsilon} < 0$

A.2 Diffusion Centrality and High Information Nodes

A.2.1 Hearing Matrix

The hearing matrix presented in Banerjee et al. (2019a) is defined as

$$\mathbf{H}(\mathbf{D}, S) = \sum_{s=1}^{S} \mathbf{D}^s \tag{16}$$

and network gossip, or the expected number of times a node j will hear a given piece of news as a function of the node of origin of the information, is

$$NG(\mathbf{D}, S)_j = H(\mathbf{D}, S)_j. \tag{17}$$

However, at single remove from the node (i.e., $S = 1$) the authors show that diffusion centrality is proportional to (out-degree centrality) (Banerjee et al., 2019a). Therefore, if we assume the caller cannot call out of their network neighborhood, then they will aim to call those with the highest centrality in their neighborhood.

---

[34]The two others in Björkegren (2019) are as follows. First, the amount of information learned in a call maps to duration. There is an one-to-one mapping of duration to information with an analytic solution, $\epsilon(d^*)$. Second, relationships with higher information flows provide more utility. The optimized utility is increasing in the optimal duration, $\frac{\partial v(d, \epsilon(d))}{\partial d} > 0$. These are both important and useful as the paper estimates a structural model to do welfare analysis. The first recovers the error or shock term from the model and the second allows for welfare analysis. However, in the context of our analysis, neither exercise is planned.

A.2.2 Refinements to Diffusion Centrality

King (2020) makes the argument that diffusion centrality double counts information flows presents three additional concepts that refine diffusion centrality: *word-of-mouth*, *obstructed*, and *visibility centrality*. The author designs word-of-mouth centrality to create a similar measure while removing the double counting. Obstruction centrality and visibility centrality both look at a case where some node does not pass information within word-of-mouth centrality. In this case, these may benefit the measurement in our context.[35] Obstructed centrality measures the average probability a node will receive signals sent by others, when that signal is obstructed by any of the other nodes in the network. The cost of this refinement from concepts of word-of-mouth centrality is that the centrality of each alter j will be defined separately for each ego i. We can restrict this to $j \in N(i)$ to be more efficient, building a sparse matrix of centralities.

Bramoullé and Genicot (2018) also presents two concepts in *targeting centrality* and *reachability*. Adapting diffusion centrality to targeted requests for favors, the authors arrive at targeting centrality. Second, they provide a formula to compute reachability of an agent in a network, or the probability a request will reach them. Again, while these measures are related to what we want, they emphasize how easy it is to target or reach a given node from the perspective of the node trying to inform them of a request. Instead we want a measure that gives a general sense of how informed we expect a node is, given their social network.

B Empirical Appendix

B.1 Variables for Heterogeneity Analysis

B.1.1 Defining Strong Ties in the Data

Since the histogram of the total duration of calls (Figure 7) does not betray any "natural" breaks in the distribution that would serve as logical points of differentiation, we opt for a simple statistical approach to imposing these cutoffs. We use the mean total call duration between individuals to differentiate strong ties from other edges in the network. Given that this duration distribution is skewed, the mean call duration is around the 81st percentile. We designate dyads with total call duration higher than this percentile to be strong ties. The Lorenz curve of total duration in the right panel of Figure 7 provides a visual depiction of this variation and associated implications for our strong tie distinction: dyads in the lowest 80% of total call duration amount to only 16% of the total call duration whereas the bottom 90% account for 27% of the dyadic call duration, or an additional 11%. We can therefore think of the edges in this decile as a reasonable grouping of "average" connections to be included or excluded from strong ties. We opt to leave them in as a measure of average ties or stronger, and choose 80% as our main strong tie cutoff. At a later stage, we will check the sensitivity of our estimates to this cutoff assumption.[36]

---

[35] In particular, we would prefer a world where the ego node is excluded from the diffusion process at hand, since information passed from the ego to an alter and back to the ego will not be informative.

[36] We considered two other alternatives, which are based on finding natural non-linearities at which to make cutoffs. A first approach is to compute the Jaccard index for each edge in the baseline network and use this to guide where to make the cut-off in total duration, which measures the proportion of friends two people hold in common:

$$\text{Jaccard}(i, j) = \frac{\text{Total number of overlapping connections between } i \text{ and } j}{\text{Total number neighbors of } i \text{ of } j} \tag{18}$$

Closely related, the second approach is to simplify this to check whether edges are "supported" or not, that is, if there exist any common connection, or if $\text{Jaccard}(i, j) > 0$. Both of these approaches appeal to the concept of *bonding social capital* , which is often measured by the triadic structure of networks. Support in particular has been shown to be important in the formation of favor networks (Jackson et al., 2012) and the communications networks of migrants (Blumenstock et al., 2019). Computation of the Jaccard index turns out to be computationally very slow in our baseline network, but we hope to do so in the future in order
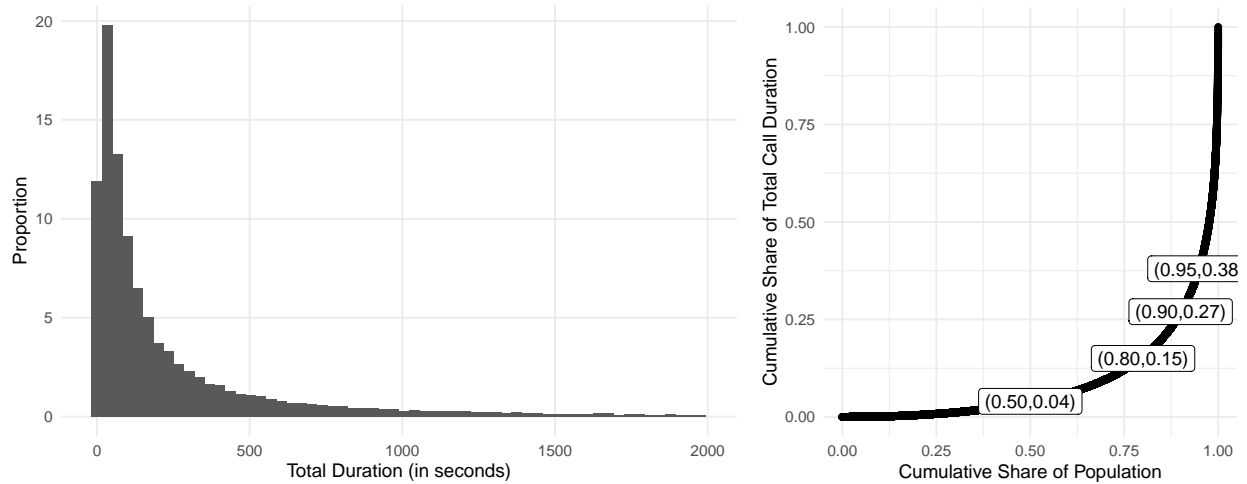
Figure 7: Histogram (left) and Lorenz curve (right) of total call duration over dyads in baseline network

B.1.2 Locating Likely Informed Nodes

We hypothesize that nodes with high degree at baseline are likely to be well informed when social unrest takes place. To proxy for likely informed nodes, we plot the degree (as well as other network statistics) of these nodes in the baseline in Figure 8. The distribution is right skewed for degree, with a median of 11, but a mean of 18. We choose a value above the mean as the cutoff for informed nodes, in this case, 30. This corresponds to roughly the 67th percentile of this statistic in the baseline.

B.2 Effects on Total Calls

Effects on total calls are presented in Figure 7. Using DID^M we estimate episodes of social unrest reduce calls by 0.103.

---

to validate the approach that we opted for.
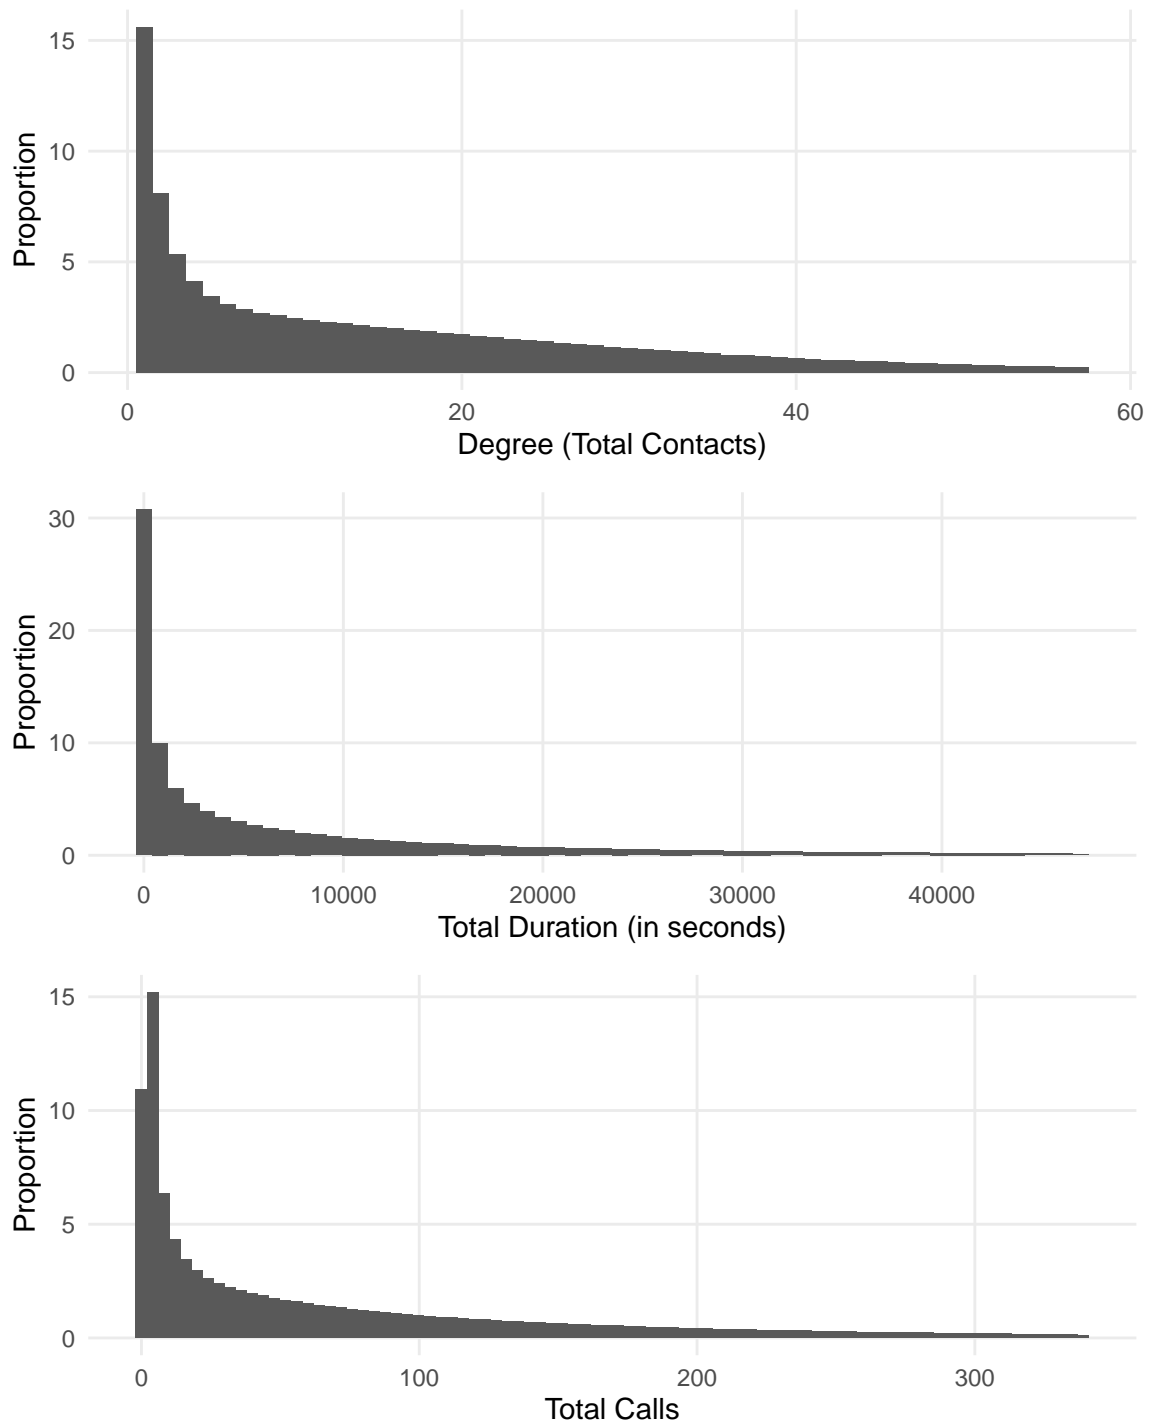
Figure 8: Histogram of node level outcomes in baseline network

Table 7: Effects of Social Unrest on Total Calls

	Total Calls		
	DiD	TWFE	DID^M
Episode of Non-Protest Social Unrest	0.031	−0.153***	−0.103***
	(0.054)	(0.050)	(0.024)
Ever Non-Protest Social Unrest	0.471***		
	(0.173)		
Day FE	Yes	Yes	
Indv. FE	No	Yes	
Cluster	Tower	Tower	Tower†
$N \times T$	3,067,920	3,067,920	3,067,920
R^2	0.005	0.617	
Adjusted R^2	0.005	0.606	
Residual Std. Error	9.753	6.141	
df RSE	3,067,882	2,982,664	
N effect			1,065,149
N switchers effect			181,124

Notes: ***Significant at the 1 percent level.
               **Significant at the 5 percent level.
               *Significant at the 10 percent level.
             † Block bootstrap with 500 repetitions

C Additional Figures

Figure 9: Protests in Port-au-Prince, January and February 2019.
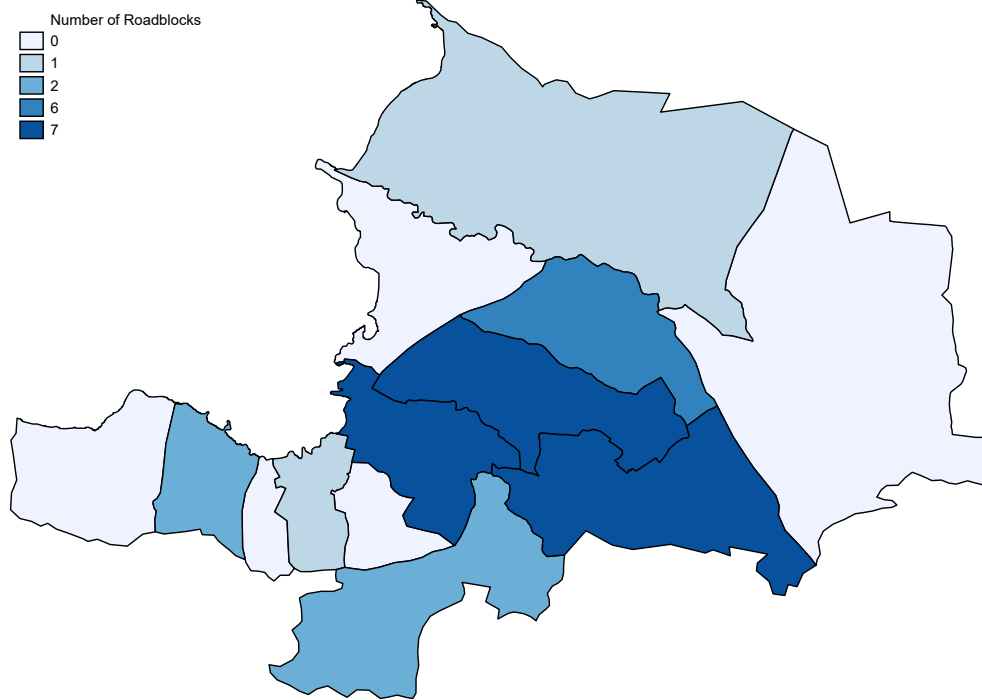
Figure 10: Roadblocks in Port-au-Prince, January and February 2019.
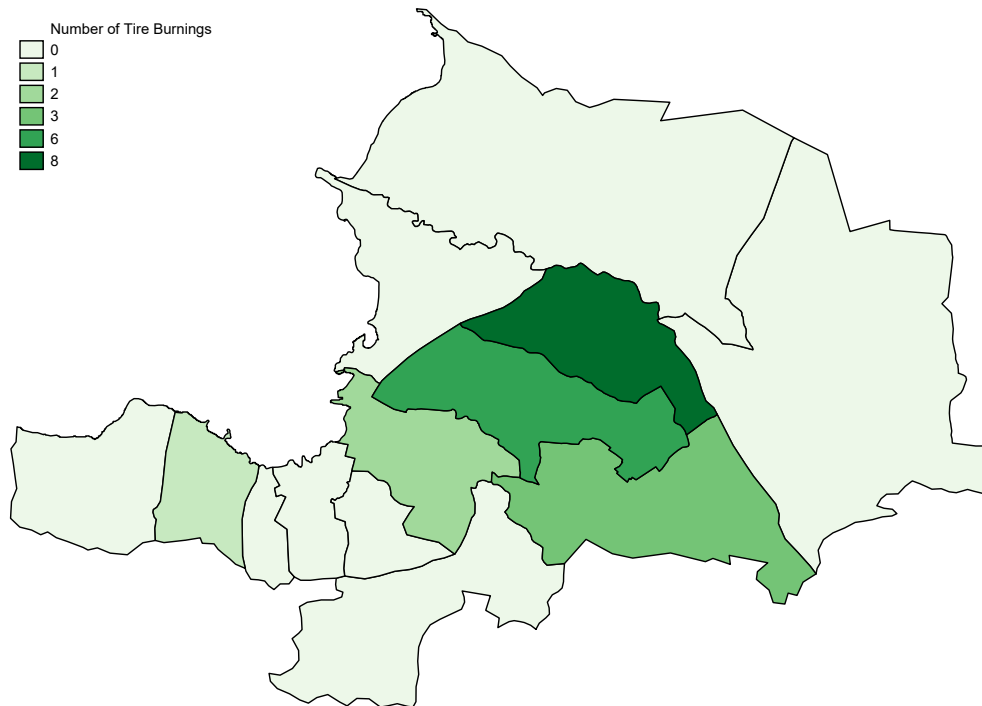
Figure 11: Tire Burning Events in Port-au-Prince, January and February 2019.
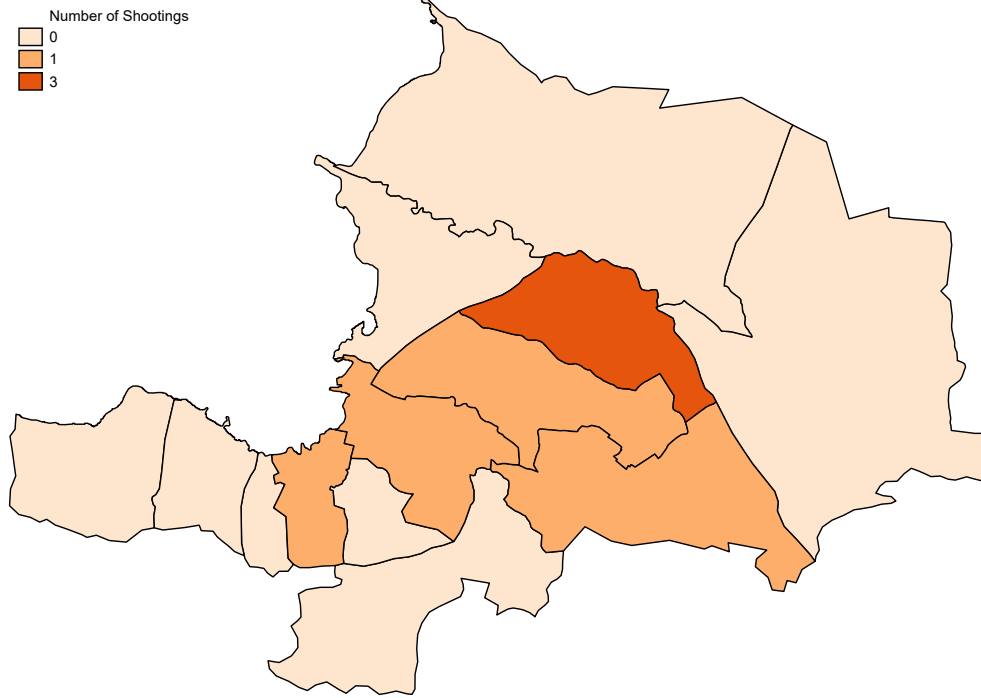
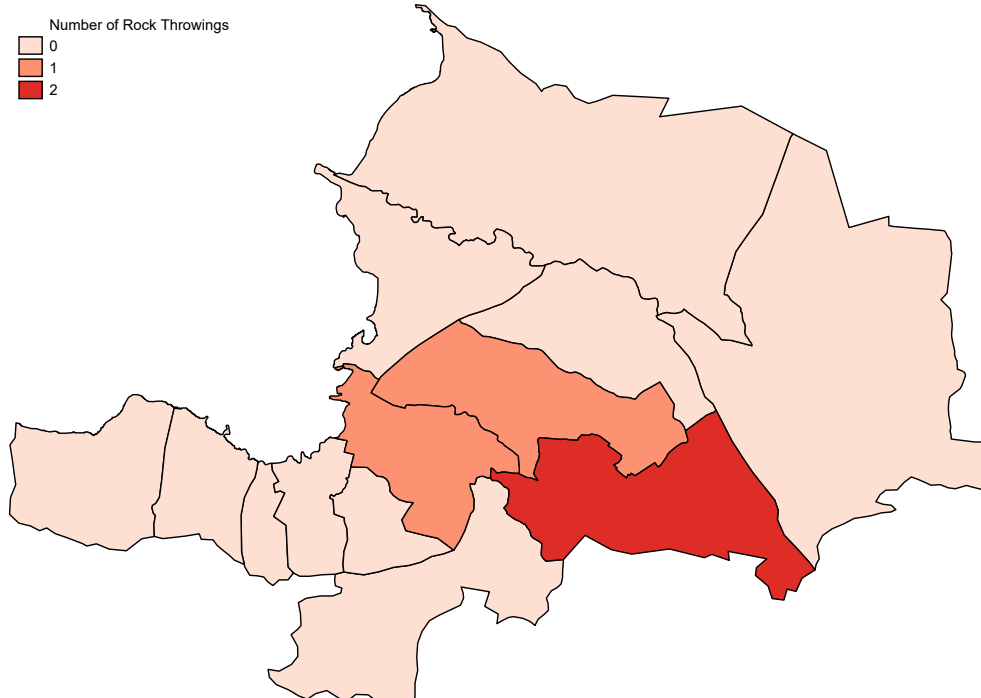Figure 12: Shootings in Port-au-Prince, January and February 2019.

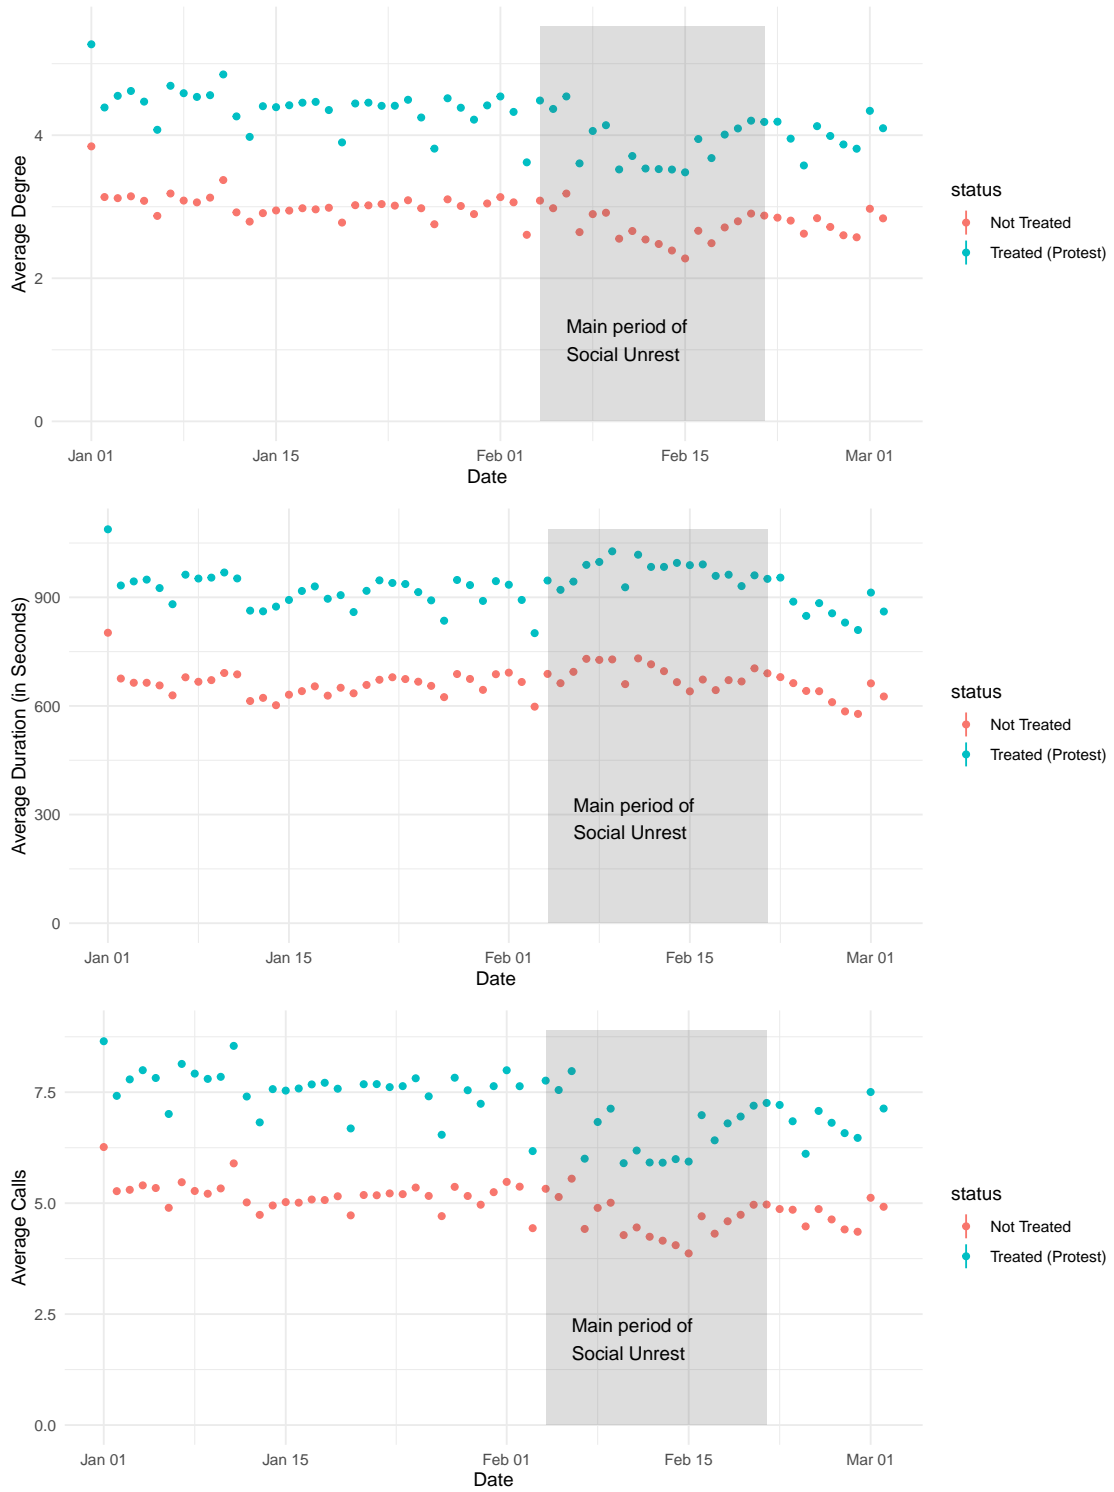Figure 13: Rock Throwing Events in Port-au-Prince, January and February 2019.

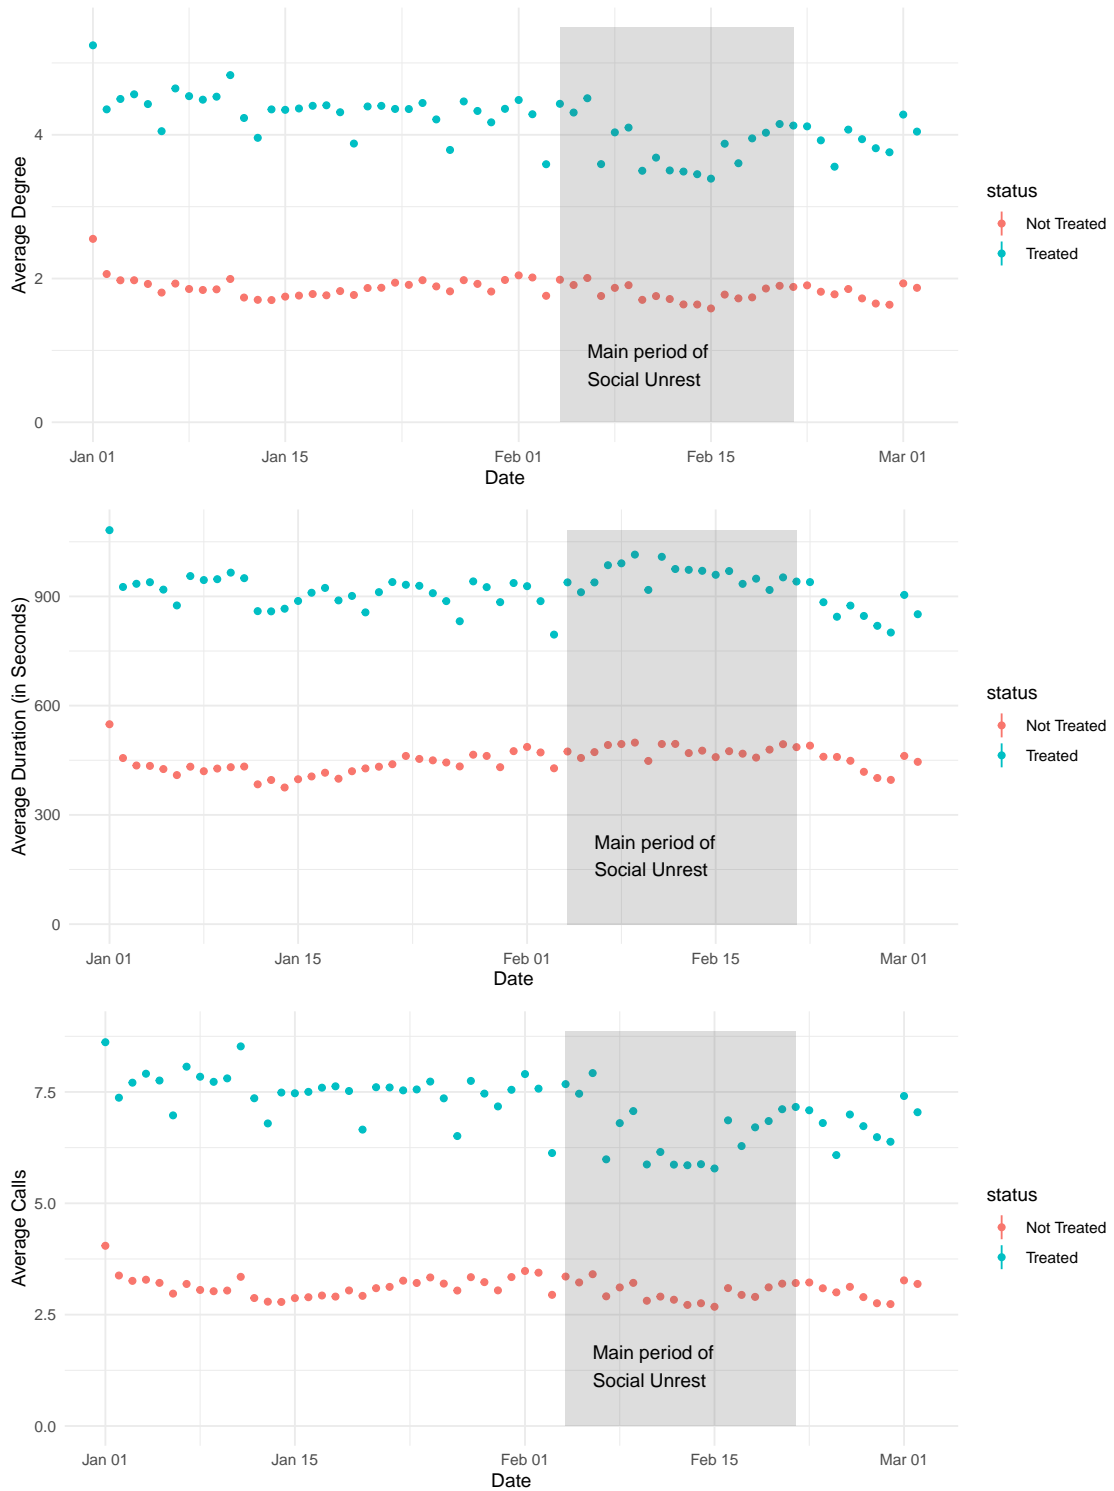Figure 14: Average degree, duration, and calls over the sample period by exposure to protest.

Figure 15: Average degree, duration, and calls over the sample period by exposure to any form of social unrest.
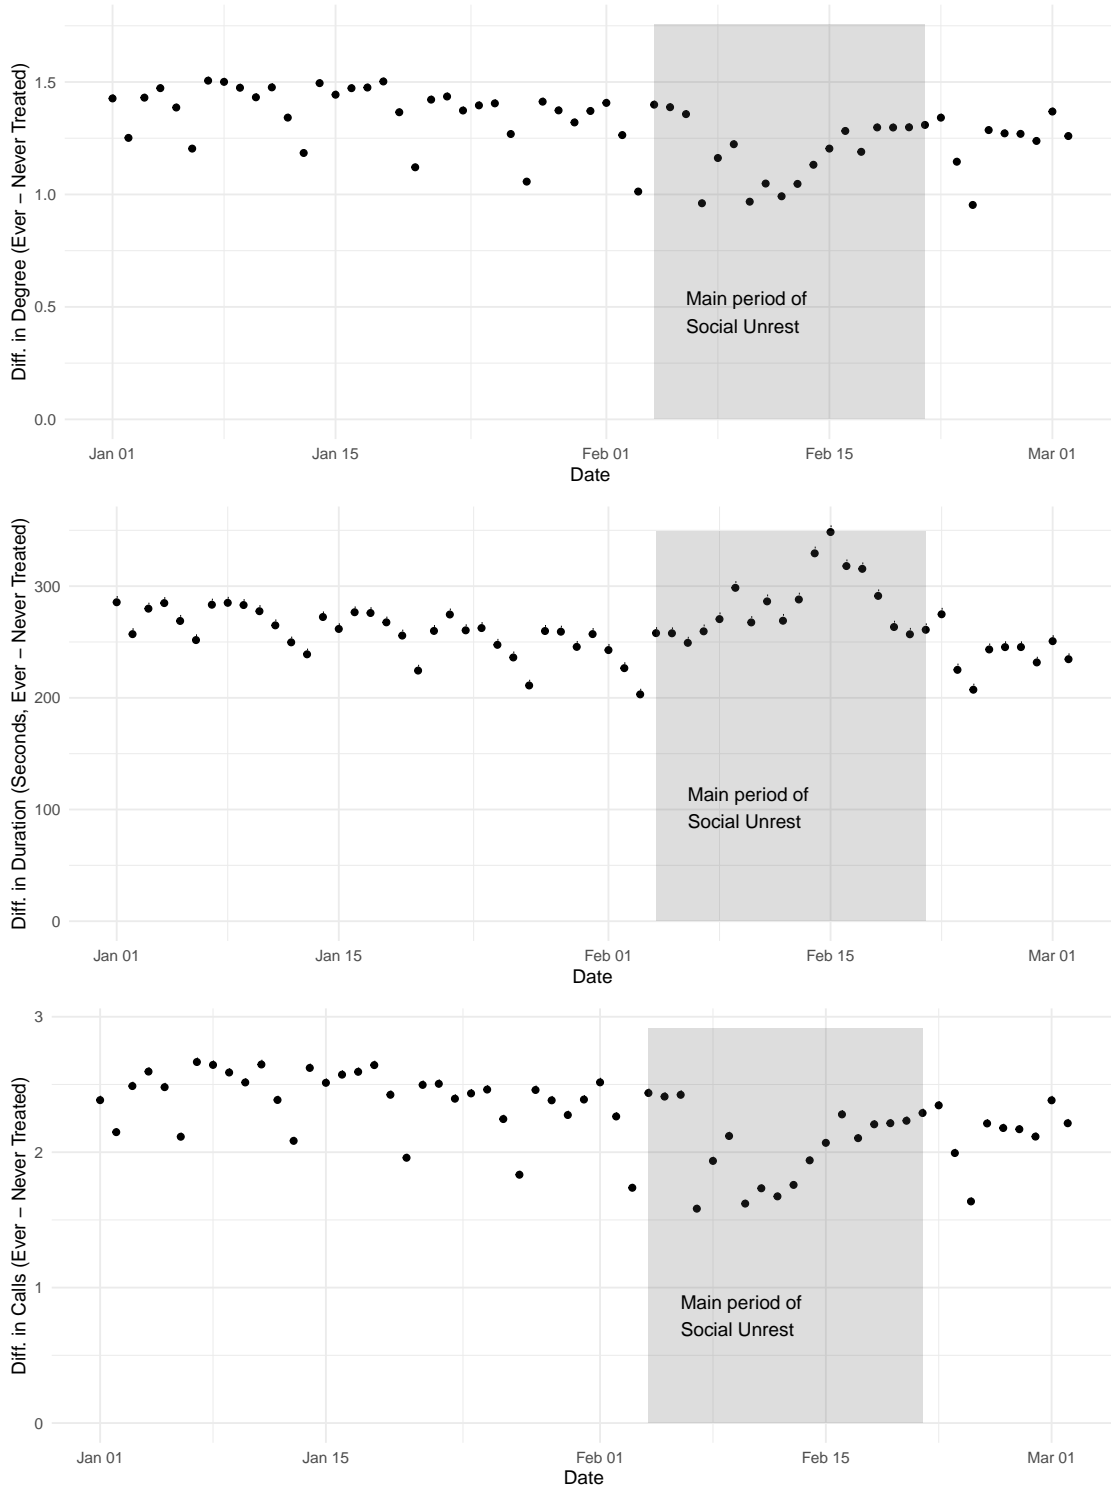
Figure 16: Difference in degree, duration, and calls over the sample period by exposure to protest.
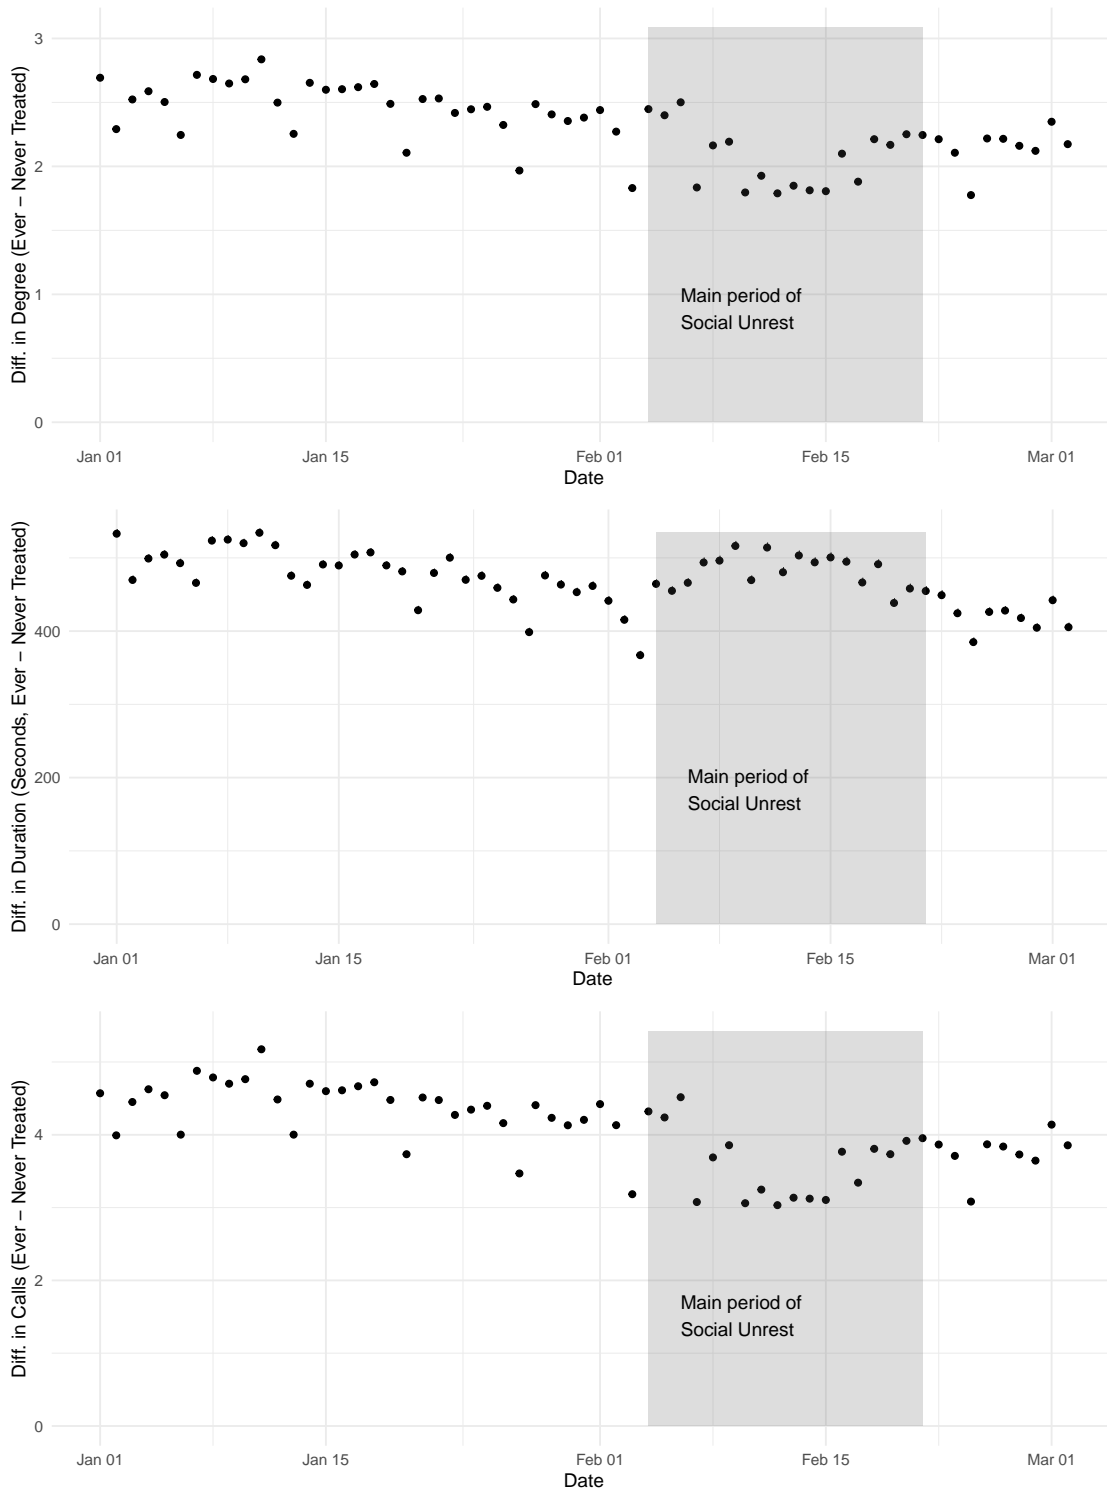
Figure 17: Difference in degree, duration, and calls over the sample period by exposure to any form of social unrest.